

Assembly Language Programming
Lecture Notes
Volume 2 – Advanced Topics

Delivered by
Belal Hashmi

Compiled by
Junaid Haroon

Table of Contents

Table of Contents	3
1 Software Interrupts	1
1.1. Interrupts	1
1.2. Hooking an Interrupt	4
1.3. BIOS and DOS Interrupts	5
2 Real Time Interrupts and Hardware Interfacing	11
2.1. Hardware Interrupts	11
2.2. I/O Ports	12
2.3. Terminate and Stay Resident	17
2.4. Programmable Interval Timer	20
2.5. Parallel Port	22
3 Debug Interrupts	31
3.1. Debugger using single step interrupt	31
3.2. Debugger using breakpoint interrupt	34
4 Multitasking	37
4.1. Concepts of Multitasking	37
4.2. Elaborate Multitasking	39
4.3. Multitasking Kernel as TSR	41
5 Video Services	47
5.1. BIOS Video Services	47
5.2. DOS Video Services	50
6 Secondary Storage	53
6.1. Physical Formation	53
6.2. Storage Access Using BIOS	54
6.3. Storage Access using DOS	59
6.4. Device Drivers	64
7 Serial Port Programming	69
7.1. Introduction	69
7.2. Serial Communication	71
8 Protected Mode Programming	73
8.1. Introduction	73
8.2. 32bit Programming	76
8.3. VESA Linear Frame Buffer	78
8.4. Interrupt Handling	80
9 Interfacing with High Level Languages	85
9.1. Calling Conventions	85
9.2. Calling C from Assembly	85
9.3. Calling Assembly from C	87

10 Comparison with Other Processors**89**

10.1. Motorola 68K Processors

89

10.2. Sun SPARC Processor

90

Software Interrupts

1.1. INTERRUPTS

Interrupts in reality are events that occurred outside the processor and the processor must be informed about them. Interrupts are asynchronous and unpredictable. Asynchronous means that the interrupts occur, independent of the working of the processor, i.e. independent of the instruction currently executing. Synchronous events are those that occur side by side with another activity. Interrupts must be asynchronous as they are generated by the external world which is unaware of the happenings inside the processor. True interrupts that occur in real time are asynchronous with the execution. Also it is unpredictable at which time an interrupt will come. The two concepts of being unpredictable and asynchronous are overlapping. Unpredictable means the time at which an interrupt will come cannot be predicted, while asynchronous means that the interrupt has nothing to do with the currently executing instruction and the current state of the processor.

The 8088 processor divides interrupts into two classes. Software interrupts and hardware interrupts. Hardware interrupts are the real interrupts generated by the external world as discussed above. Software interrupts on the contrary are not generated from outside the processor. They just provide an extended far call mechanism. Far call allows us to jump anywhere in the whole megabyte of memory. To return from the target we place both the segment and offset on the stack. Software interrupts show a similar behavior. It however pushes one more thing before both the segment and offset and that is the FLAGS register. Just like the far call loads new values in CS and IP, the interrupt call loads new values in CS, IP, and FLAGS. Therefore the only way to retain the value of original FLAGS register is to push and pop as part of interrupt call and return instructions. Pushing and popping inside the routine will not work as the routine started with an already tampered value.

The discussion of real time interrupts is deferred till the next chapter. They play the critical part in control applications where external hardware must be control and events and changes therein must be appropriately responded by the processor. To generate an interrupt the INT instruction is used. The routine that executes in response to an INT instruction is called the interrupt service routine (ISR) or the interrupt handler. Taking example from real time interrupts the routine to instruct an external hardware to close the valve of a boiler in response to an interrupt from the pressure sensor is an interrupt routine.

The software interrupt mechanism in 8088 uses vectored interrupts meaning that the address of the interrupt routine is not directly mentioned in an interrupt call, rather the address is looked up from a table. 8088 provides a mechanism for mapping interrupts to interrupt handlers. Introducing a new entry in this mapping table is called hooking an interrupt.

Syntax of the INT instruction is very simple. It takes a single byte argument varying from 0-255. This is the interrupt number informing the processor, which interrupt is currently of interest. This number correlates to the interrupt handler routine by a routing or vectoring mechanism. A few interrupt numbers in the start are reserved and we generally do not use them. They are related to the processor working. For example INT 0 is the

divide by zero interrupt. A list of all reserved interrupts is given later. Such interrupts are programmed in the hardware to generate the designated interrupt when the specified condition arises. The remaining interrupts are provided by the processor for our use. Some of these were reserved by the IBM PC designers to interface user programs with system software like DOS and BIOS. This was the logical choice for them as interrupts provided a very flexible architecture. The remaining interrupts are totally free for use in user software.

The correlation process from the interrupt number to the interrupt handler uses a table called interrupt vector table. Its location is fixed to physical memory address zero. Each entry of the table is four bytes long containing the segment and offset of the interrupt routine for the corresponding interrupt number. The first two bytes in the entry contain the offset and the next two bytes contain the segment. The little endian rule of putting the more significant part (segment) at a higher address is seen here as well. Mathematically offset of the interrupt n will be at $nx4$ while the segment will be at $nx4+2$. One entry in this table is called a vector. If the vector is changed for interrupt 0 then INT 0 will take execution to the new handler whose address is now placed at those four bytes. INT 1 vector occupies location 4, 5, 6, and 7 and similarly vector for INT 2 occupies locations 8, 9, 10, and 11. As the table is located in RAM it can be changed anytime. Immediately after changing it the interrupt mapping is changed and now the interrupt will result in execution of the new routine. This indirection gives the mechanism extreme flexibility.

The operation of interrupt is same whether it is the result of an INT instruction (software interrupt) or it is generated by an external hardware which passes the interrupt number by a different mechanism. The currently executing instruction is completed, the current value of FLAGS is pushed on the stack, then the current code segment is pushed, then the offset of the next instruction is pushed. After this it automatically clears the trap flag and the interrupt flag to disallow further interrupts until the current routine finishes. After this it loads the word at $nx4$ in IP and the word at $nx4+2$ in CS if interrupt n was generated. As soon as these values are loaded in CS and IP execution goes to the start of the interrupt handler. When the handler finishes its work it uses the IRET instruction to return to the caller. IRET pops IP, then CS, and then FLAGS. The original value of IF and TF is restored which reenables further interrupts. IF and TF will be discussed in detail in the discussion of real time interrupts. We have discussed three things till now.

1. The INT and IRET instruction format and syntax
2. The formation of IVT (interrupt vector table)
3. Operation of the processor when an interrupt is generated

Just as discussed in the subroutines chapter, the processor will not match interrupt calls to interrupt returns. If a RETF is used in the end of an ISR the processor will still return to the caller but the FLAGS will remain on the stack which will destroy the expectations of the caller with the stack. If we know what we are doing we may use such different combination of instructions. Generally we will use IRET to return from an interrupt routine. Apart from indirection the software interrupt mechanism is similar to CALL and RET. Indirection is the major difference.

The operation of INT can be written as:

- $sp \leftarrow sp+2$
- $[sp] \leftarrow \text{flag}$
- $sp \leftarrow sp+2$
- $if \leftarrow 0$
- $tf \leftarrow 0$
- $[sp] \leftarrow cs$
- $sp \leftarrow sp+2$
- $[sp] \leftarrow ip$

- $ip \leftarrow [0:N*4]$
- $cs \leftarrow [0:N*4+2]$

The operation of IRET can be written as:

- $ip \leftarrow [sp]$
- $sp \leftarrow sp-2$
- $cs \leftarrow [sp]$
- $sp \leftarrow sp-2$
- $flag \leftarrow [sp]$
- $sp \leftarrow sp-2$

The above is the microcode description of INT and IRET. To obey an assembly language instruction the processor breaks it down into small operations. By reading the microcode of an instruction its working can be completely understood.

The interrupt mechanism we have studied is an extended far call mechanism. It pushes FLAGS in addition to CS and IP and it loads CS and IP with a special mechanism of indirection. It is just like the table of contents that is located at a fixed position and allows going directly to chapter 3, to chapter 4 etc. If this association is changed in the table of contents the direction of the reader changes. For example if Chapter 2 starts at page 220 while 240 is written in the table of contents, the reader will go to page 240 and not 220. The table of contents entry is a vector to point to map the chapter number to page number. IVT has 256 chapters and the interrupt mechanism looks up the appropriate chapter number to reach the desired page to find the interrupt routine.

Another important similarity is that table of contents is always placed at the start of the book, a well known place. Its physical position is fixed. If some publishers put it at some place, others at another place, the reader will be unable to find the desired chapter. Similarly in 8088 the physical memory address zero is fixed for the IVT and it occupies exactly a kilobyte of memory as the $256 \times 4 = 1K$ where 256 is the number of possible interrupt vectors while the size of one vector is 4 bytes.

Interrupts introduce temporary breakage in the program flow, sometimes programmed (software interrupts) and unprogrammed at other times (hardware interrupts). By hooking interrupts various system functionalities can be controlled. The interrupts reserved by the processor and having special functions in 8088 are listed below:

- INT 0, Division by zero
Meaning the quotient did not fit in the destination register. This is a bit different as this interrupt does not return to the next instruction, rather it returns to the same instruction that generated it, a DIV instruction ofcourse. Here INT 0 is automatically generated by a DIV when a specific situation arises, there is no INT 0 instruction.
- INT 1, Trap, Single step Interrupt
This interrupt is used in debugging with the trap flag. If the trap flag is set the Single Step Interrupt is generated after every instruction. By hooking this interrupt a debugger can get control after every instruction and display the registers etc. 8088 was the first processor that has this ability to support debugging.
- INT 2, NMI-Non Maskable Interrupt
Real interrupts come from outside the processor. INT 0 is not real as it is generated from inside. For real interrupts there are two pins in the processor, the INT pin and the NMI pin. The processor can be directed to listen or not to listen to the INT pin. Consider a recording studio, when the recording is going on, doors are closed so that no interruption occurs, and when there is a break, the doors are opened so that if someone is waiting outside can come in. However if there is an urgency like fire outside then the door must be broken and the recording must not be catered for. For such situations is the NMI pin

which informs about fatal hardware failures in the system and is tied to interrupt 2. INT pin can be masked but NMI cannot be masked.

- INT 3, Debug Interrupt
The only special thing about this interrupt is that it has a single byte opcode and not a two byte combination where the second byte tells the interrupt number. This allows it to replace any instruction whatsoever. It is also used by the debugger and will be discussed in detail with the debugger working.
- INT 4, Arithmetic Overflow, change of sign bit
The overflow flag is set if the sign bit unexpectedly changes as a result of a mathematical or logical instruction. However the overflow flag signals a real overflow only if the numbers in question are treated as signed numbers. So this interrupt is not automatically generated but as a result of a special instruction INTO (interrupt on overflow) if the overflow flag is set. Otherwise the INTO instruction behaves like a NOP (no operation).

These are the five interrupts reserved by Intel and are generally not used in our operations.

1.2. HOOKING AN INTERRUPT

To hook an interrupt we change the vector corresponding to that interrupt. As soon as the interrupt vector changes, that interrupt will be routed to the new handler. Our first example is with the divide by zero interrupt. The normal system defined behavior in response to divide by zero is to display an error message and terminate the program. We will change it to display our own message.

Example 8.1

```

001      ; hooking divide by zero interrupt
002      [org 0x0100]
003              jmp  start
004
005      message:      db  'You divided something by zero.', 0
006
007-029      ;;;; COPY LINES 028-050 FROM EXAMPLE 7.4 (strlen) ;;;;
030-049      ;;;; COPY LINES 005-024 FROM EXAMPLE 7.1 (clrscr) ;;;;
050-090      ;;;; COPY LINES 050-090 FROM EXAMPLE 7.4 (printstr) ;;;;
091
092      ; divide by zero interrupt handler
093      myisrfor0:      push ax                      ; push all regs
094                      push bx
095                      push cx
096                      push dx
097                      push si
098                      push di
099                      push bp
100                      push ds
101                      push es
102
103                      push cs
104                      pop  ds                      ; point ds to our data segment
105
106                      call clrscr                  ; clear the screen
107                      mov  ax, 30
108                      push ax                      ; push x position
109                      mov  ax, 20
110                      push ax                      ; push y position
111                      mov  ax, 0x71                ; white on blue attribute
112                      push ax                      ; push attribute
113                      mov  ax, message
114                      push ax                      ; push offset of message
115                      call printstr                ; print message
116
117                      pop  es
118                      pop  ds
119                      pop  bp

```


120	pop di	
121	pop si	
123	pop dx	
124	pop cx	
125	pop bx	
126	pop ax	
127	iret	; return from interrupt
128		
129	; subroutine to generate a divide by zero interrupt	
130	genint0:	mov ax, 0x8432 ; load a big number in ax
131		mov bl, 2 ; use a very small divisor
132		div bl ; interrupt 0 will be generated
133		ret
134		
135	start:	xor ax, ax
136		mov es, ax ; load zero in es
137		mov word [es:0*4], myisrfor0 ; store offset at n*4
138		mov [es:0*4+2], cs ; store segment at n*4+2
139		call genint0 ; generate interrupt 0
140		
141		mov ax, 0x4c00 ; terminate program
142		int 0x21
93-101	We often push all registers in an interrupt service routine just to be sure that no unintentional modification to any register is made. Since any code may be interrupted an unintentional modification will be hard to debug	
103-104	Since interrupt can be called from anywhere we are not sure about the value in DS so we reset it to our code segment.	

When this program is executed our desired message will be shown instead of the default message and the computer will hang thereafter. The first thing to observe is that there is no INT 0 call anywhere in the code. INT 0 was invoked automatically by an internal mechanism of the processor as a result of the DIV instruction producing a result that cannot fit in the destination register. Just by changing the vector we have changed the response of the system to divide overflow situations.

However the system stuck instead of returning to the next instruction. This is because divide overflow is a special type of interrupt that returns to the same instruction instead of the next instruction. This is why the default handler forcefully terminates the program instead of returning. Now the IRET will take control back to the DIV instruction which will again generate the same interrupt. So the computer is stuck in an infinite loop.

1.3. BIOS AND DOS INTERRUPTS

In IBM PC there are certain interrupts designated for user programs to communicate with system software to access various standard services like access to the floppy drive, hard drive, vga, clock etc. If the programmer does not use these services he has to understand the hardware details like which particular controller is used and how it works. To avoid this and provide interoperability a software interface to basic hardware devices is provided except in very early computers. Since the manufacturer knows the hardware it burns the software to control its hardware in ROM. Such software is called firmware and access to this firmware is provided through specified interrupts.

This basic interface to the hardware is called BIOS (basic input output services). When the computer is switched on, BIOS gets the control at a specified address. The messages at boot time on the screen giving BIOS version, detecting different hardware are from this code. BIOS has the responsibility of testing the basic hardware including video, keyboard, floppy drive, hard drive etc and a special program to bootstrap. Bootstrap means to load OS from hard disk and from there OS takes control and proceeds to load its components and display a command prompt in the end. There are two

important programs; BIOS and OS. OS services are high level and build upon the BIOS services. BIOS services are very low level. A level further lower is only directly controlling the hardware. BIOS services provide a hardware independent layer above the hardware and OS services provide another higher level layer over the BIOS services. We have practiced direct hardware access with the video device directly without using BIOS or DOS. The layer of BIOS provides services like display a character, clear the screen, etc. All these layers are optional in that we can skip to whatever lower layer we want.

The most logical way to provide access to firmware is to use the interrupt mechanism. Specific services are provided at specific interrupts. CALL could also have been used but in that case every manufacturer would be required to place specific routines at specific addresses, which is not a flexible mechanism. Interrupts provide standard interrupt number for the caller and flexibility to place the interrupt routine anywhere in the memory for the manufacturer. Now for the programmer it is decided that video services will be provided at INT 10 but the actual address of the video services can and do vary on computers from different manufacturers. Any computer that is IBM compatible must make the video services accessible through INT 10. Similarly keyboard services are available at INT 16 and this is standard in every IBM compatible. Manufacturers place the code wherever they want and the services are exported through this interrupt.

BIOS exports its various services through different interrupts. Keyboard services are exported through INT 16, parallel port services through INT 17 and similarly others through different interrupts. DOS has a single entry point through INT 21 just like a pin hole camera, this single entry point leads to a number of DOS services. So how one interrupt provides a number of different services. A concept of service number is used here which is a de facto standard in providing multiple services through an interrupt. INT 10 is for video services and each of character printing service, screen clearing service, cursor movement service etc. has a service number associated to it. So we say INT 10 service 0 is used for this purpose and INT 10 service 1 is used for that purpose etc. Service numbers for different standard services are also fixed for every IBM compatible. The concept of exported services through interrupts is expanded with the service numbering scheme.

The service number is usually given in the AH register. Sometimes these 256 services seem less. For example DOS exports thousands of services. So we will often see an extension to a level further with subservices. For example INT 10 character generator services are all provided through a single service number and the services are distinguished with a subservice number.

The finally selected service would need some arguments for it to work. In interrupts arguments are usually not given through stack, rather registers are used. The BIOS and DOS specifications list which register contains which argument for a particular service of a particular interrupt.

We will touch some important BIOS and DOS services and not cover it completely neither is it possible to cover it in this space. A very comprehensive reference of interrupts is the Ralph Brown List. It is just a reference and not to be studied from end to end. All interrupts cannot be remembered and there is no need to remember them.

The service number is almost always in AH while the subservice number is in AL or BL and sometimes in other registers. The documentation of the service we are using will list which register should hold what when the interrupt is invoked for that particular service.

Our first target using BIOS is video so let us proceed to our first program that uses INT 10 service 13 to print a string on the screen. BIOS will work even if the video memory is not at B8000 (a very old video card) since BIOS knows everything about the hardware and is hardware specific.

Example 8.2	
001	; print string using bios service
002	[org 0x0100]
003	jmp start
004	message: db 'Hello World'
005	
006	start: mov ah, 0x13 ; service 13 - print string
007	mov al, 1 ; subservice 01 - update cursor
008	mov bh, 0 ; output on page 0
009	mov bl, 7 ; normal attrib
010	mov dx, 0x0A03 ; row 10 column 3
011	mov cx, 11 ; length of string
012	push cs
013	pop es ; segment of string
014	mov bp, message ; offset of string
015	int 0x10 ; call BIOS video service
016	
017	mov ax, 0x4c00 ; terminate program
018	int 0x21
007	The subservice are versions of printstring that update and do not update the cursor after printing the string etc.
008	Text video screen is in the form of pages which can be upto 32. At one time one page is visible which is by default the zeroth page unless we change it.

When we execute it the string is printed and the cursor is updated as well. With direct access to video memory we had no control over the cursor. To control cursor a different mechanism to access the hardware was needed.

Our next example uses the keyboard service to read a key. The combination of keyboard and video services is used in almost every program that we see and use. We will wait for four key presses; clear the screen after the first, and draw different strings after the next key presses and exiting after the last. We will use INT 16 service 1 for this purpose. This is a blocking service so it does not return until a key has been pressed. We also used the blinking attribute in this example.

Example 8.3	
001	; print string and keyboard wait using BIOS services
002	[org 0x100]
003	jmp start
004	
005	msg1: db 'hello world', 0
006	msg2: db 'hello world again', 0
007	msg3: db 'hello world again and again', 0
008	
009-028	;;;; COPY LINES 005-024 FROM EXAMPLE 7.1 (clrscr) ;;;;
029-069	;;;; COPY LINES 050-090 FROM EXAMPLE 7.4 (printstr) ;;;;
070-092	;;;; COPY LINES 028-050 FROM EXAMPLE 7.4 (strlen) ;;;;
093	
094	start: mov ah, 0x10 ; service 10 - vga attributes
095	mov al, 03 ; subservice 3 - toggle blinking
096	mov bl, 01 ; enable blinking bit
097	int 0x10 ; call BIOS video service
098	
099	mov ah, 0 ; service 0 - get keystroke
100	int 0x16 ; call BIOS keyboard service
101	
102	call clrscr ; clear the screen
103	
104	mov ah, 0 ; service 0 - get keystroke
105	int 0x16 ; call BIOS keyboard service
106	
107	mov ax, 0
108	push ax ; push x position
109	mov ax, 0
110	push ax ; push y position
111	mov ax, 1 ; blue on black

112	push ax	; push attribute
113	mov ax, msg1	
114	push ax	; push offset of string
115	call printstr	; print the string
116		
117	mov ah, 0	; service 0 - get keystroke
118	int 0x16	; call BIOS keyboard service
119		
120	mov ax, 0	
121	push ax	; push x position
123	mov ax, 0	
124	push ax	; push y position
125	mov ax, 0x71	; blue on white
126	push ax	; push attribute
127	mov ax, msg1	
128	push ax	; push offset of string
129	call printstr	; print the string
130		
131	mov ah, 0	; service 0 - get keystroke
132	int 0x16	; call BIOS keyboard service
133		
134	mov ax, 0	
135	push ax	; push x position
136	mov ax, 0	
137	push ax	; push y position
138	mov ax, 0xF4	; red on white blinking
139	push ax	; push attribute
140	mov ax, msg1	
141	push ax	; push offset of string
142	call printstr	; print the string
143		
144	mov ah, 0	; service 0 - get keystroke
145	int 0x16	; call BIOS keyboard service
146		
147	mov ax, 0x4c00	; terminate program
148	int 0x21	
099-100	This service has no parameters so only the service number is initialized in AH. This is the only service so there is no subservice number as well. The ASCII code of the char pressed is returned in AL after this service.	

EXERCISES

1. Write a TSR that forces a program to exit when it tries to become a TSR using INT 21h/Service 31h by converting its call into INT 21h/Service 4Ch.
2. Write a function to clear the screen whose only parameter is always zero. The function is hooked at interrupt 80h and may also be called directly both as a near call and as a far call. The function should detect how it is called and return appropriately. It is provided that the direction flag will be set before the function is called.
3. Write a function that takes three parameters, the interrupt number (N) and the segment and offset of an interrupt handler XISR. The arguments are pushed in the order N, XISR's offset and XISR's segment. It is known that the first two instructions of XISR are PUSHF and CALL 0:0 followed by the rest of the interrupt handler. PUSHF instruction is of one byte and far call is of 5 bytes with the first byte being the op-code, the next two containing the target offset and the last two containing the target segment. The function should hook XISR at interrupt N and chain it to the interrupt handler previously hooked at N by manipulating the call 0:0 instruction placed near the start of XISR.
4. Write a TSR that provide the circular queue services via interrupt 0x80 using the code written in Exercise 5.XX. The interrupt procedure should call one of qcreate, qdestroy, qempty, qadd,

qremove, and uninstall based on the value in AH. The uninstall function should restore the old interrupt 0x80 handler and remove the TSR from memory.

Real Time Interrupts and Hardware Interfacing

2.1. HARDWARE INTERRUPTS

The same mechanism as discussed in the previous chapter is used for real interrupts that are generated by external hardware. However there is a single pin outside the processor called the INT pin that is used by external hardware to generate interrupts. The detailed operation that happens outside the process when an interrupt is generated is complex and only a simplified view will be discussed here; the view that is relevant to an assembly language programmer. There are many external devices that need the processor's attention like the keyboard, hard disk, floppy disk, sound card. All of them need real time interrupts at some point in their operation. For example if a program is busy in some calculations for three minutes the key strokes that are hit meanwhile should not be wasted. Therefore when a key is pressed, the INT signal is sent, an interrupt generated and the interrupt handler stores the key for later use. Similarly when the printer is busy printing we cannot send it more data. As soon as it gets free from the previous job it interrupts the processor to inform that it is free now. There are many other examples where the processor needs to be informed of an external event. If the processor actively monitors all devices instead of being automatically interrupted then it there won't be any time to do meaningful work.

Since there are many devices generating interrupts and there is only one pin going inside the processor and one pin cannot be technically driven by more than one source a controller is used in between called the Programmable Interrupt Controller (PIC). It has eight input signals and one output signal. It assigns priorities to its eight input pins from 0 to 7 so that if more than one interrupt comes at the same times, the highest priority one is forwarded and the rest are held till that is serviced. The rest are forwarded one by one according to priority after the highest priority one is completed. The original IBM XT computer had one PIC so there were 8 possible interrupt sources. However IBM AT and later computers have two PIC totaling 16 possible interrupt sources. They are arranged in a special cascade master slave arrangement so that only one output signal comes towards the processor. However we will concentrate on the first interrupt controller only.

The priority can be understood with the following example. Consider eight parallel switches which are all closed and connected to form the output signal. When a signal comes on one of the switches, it is passed on to the output and this switch and all below it are opened so that no further signals can pass through it. The higher priority switches are still closed and the signal on them can be forwarded. When the processor signals that it is finished with the processing the switches are closed again and any waiting interrupts may be forwarded. The way the processor signals ending of the interrupt service routine is by using a special mechanism discussed later.

The eight input signals to the PIC are called Interrupt Requests (IRQ). The eight lines are called IRQ 0 to IRQ 7. These are the input lines of the 8451.* For example IRQ 0 is driven by a timer device. The timer device keeps

* 8451 is the technical number of the PIC.

generating interrupts with a specified frequency. IRQ 1 is driven by the keyboard when generates an interrupts when a key is pressed or released. IRQ 2 is the cascading interrupt connected to the output of the second 8451 in the machine. IRQ 3 is connected to serial port COM 2 while IRQ 4 is connected to serial port COM 1. IRQ 5 is used by the sound card or the network card or the modem. An IRQ conflict means that two devices in the system want to use the same IRQ line. IRQ 6 is used by the floppy disk drive while IRQ 7 is used by the parallel port.

Each IRQ is mapped to a specific interrupt in the system. This is called the IRQ to INT mapping. IRQ 0 to IRQ 7 are consecutively mapped on interrupts 8 to F. This mapping is done by the PIC and not the processor. The actual mechanism fetches one instruction from the PIC whenever the INT pin is signaled instead of the memory. We can program the PIC to generate a different set of interrupts on the same interrupt requests. From the perspective of an assembly language programmer an IRQ 0 is translated into an INT 8 without any such instruction in the program and that's all. Therefore an IRQ 0, the highest priority interrupt, is generated by the timer chip at a precise frequency and the handler at INT 8 is invoked which updates the system time. A key press generates IRQ 1 and the INT 9 handler is invoked which stores this key. To handler the timer and keyboard interrupts one can replace the vectors corresponding to interrupt 8 and 9 respectively. For example if the timer interrupt is replaced and the floppy is accessed by some program, the floppy motor and its light will remain on for ever as in the normal case it is turned off by the timer interrupt after two seconds in anticipation that another floppy access might be needed otherwise the time motor takes to speed up will be needed again.[†]

We have seen that an interrupt request from a device enters the PIC as an IRQ, from there it reaches the INT pin of the processor, the proessor receives the interrupt number from the PIC, generates the designated interrupt, and finally the interrupt handler gain control and can do whatever is desired. At the end of servicing the interrupt the handler should inform the PIC that it is completed so that lower priority interrupts can be sent from the PIC. This signal is called an End Of Interrupt (EOI) signal and is sent through the I/O ports of the interrupt controller.

2.2. I/O PORTS

There are hundreds of peripheral devices in the system, PIC is one example. The processor needs to communicate with them, give and take data from them, otherwise their presence is meaningless. Memory has a totally different purpose. It contains the program to be executed and its data. It does not control any hardware. For communicating with peripheral devices the processor uses I/O ports. There are only two operations with the external world possible, read or write. Similarly with I/O ports the processor can read or write an I/O port. When an I/O port is read or written to, the operation is not as simple as it happens in memory. Some hardware changes it functionality or performs some operation as a result.

IBM PC has separate memory address space and peripheral address space. Some processors use memory mapped I/O in which case designated memory cells work as ports for specific devices. In case of Intel a special pin on the control bus signals whether the current read or write is from the memory address space or from the peripheral address space. The same address and data buses are used to select a port and to read or write data from that port. However with I/O only the lower 16 bits of the address bus are used meaning that there are a total of 65536 possible I/O ports. Now keyboard has special

[†] The programs discussed from now onwards in the book must be executed in pure DOS and not in a DOS window so that we are in total control of the PIC and other devices.

I/O ports designated to it, PIC has others, DMA, sound card, network card, each has some ports.

If the two address spaces are differentiated in hardware, they must also have special instructions to select the other address space. We have the IN and OUT instructions to read or write from the peripheral address space. When MOV is given the processor selects the memory address space, when IN is given the processor selects the peripheral address space.

IN and OUT instructions

The IN and OUT instructions have a byte form and a word form but the byte form is almost always used. The source register in OUT and destination register in IN is AL or AX depending on which form is used. The port number can be directly given in the instruction if it fits in a byte otherwise it has to be given in the DX register. Port numbers for specific devices are fixed by the IBM standard. For example 20 and 21 are for PIC, 60 to 64 for Keyboard, 378 for the parallel port etc. A few examples of IN and OUT are below:

```
in al, 0x21
mov dx, 0x378
in al, dx
out 0x21, al
mov dx, 0x378
out dx, al
```

PIC Ports

Programmable interrupt controller has two ports 20 and 21. Port 20 is the control port while port 21 is the interrupt mask register which can be used for selectively enabling or disabling interrupts. Each of the bits at port 21 corresponds to one of the IRQ lines. We first write a small program to disable the keyboard using this port. As we know that the keyboard IRQ is 1, we place a 1 bit at its corresponding position. A 0 bit will enable an interrupt and a 1 bit disables it. As soon as we write it on the port keyboard interrupts will stop arriving and the keyboard will effectively be disabled. Even Ctrl-Alt-Del would not work; the reset power button has to be used.

Example 9.1	
001	; disable keyboard interrupt in PIC mask register
002	[org 0x0100]
003	in al, 0x21 ; read interrupt mask register
004	or al, 2 ; set bit for IRQ2
005	out 0x21, al ; write back mask register
006	
007	mov ax, 0x4c00 ; terminate program
008	int 0x21

After this three line mini program is executed the computer will not understand anything else. Its ears are closed. No keystrokes are making their way to the processor. Ports always make something happen on the system. A properly designed system can launch a missile on writing a bit on some port. Memory is simple in that it is all that it is. In ports every bit has a meaning that changes something in the system.

As we previously discussed every interrupt handler invoked because of an IRQ must signal an EOI otherwise lower priority interrupts will remain disabled.

Keyboard Controller

We will go in further details of the keyboard and its relation to the computer. We will not discuss how the keyboard communicates with the keyboard controller in the computer rather we will discuss how the keyboard

controller communicates with the processor. Keyboard is a collection of labeled buttons and every button is designated a number (not the ASCII code). This number is sent to the processor whenever the key is pressed. From this number called the scan code the processor understands which key was pressed. For each key the scan code comes twice, once for the key press and once for the key release. Both are scan codes and differ in one bit only. The lower seven bits contain the key number while the most significant bit is clear in the press code and set in the release code. The IBM PC standard gives a table of the scan codes of all keys.

If we press Shift-A resulting in a capital A on the screen, the controller has sent the press code of Shift, the press code of A, the release code of A, the release code of Shift and the interrupt handler has understood that this sequence should result in the ASCII code of 'A'. The 'A' key always produces the same scan code whether or not shift is pressed. It is the interrupt handler's job to remember that the press code of Shift has come and release code has not yet come and therefore to change the meaning of the following key presses. Even the capslock key works the same way.

An interesting thing is that the two shift keys on the left and right side of the keyboard produce different scan codes. The standard way implemented in BIOS is to treat that similarly. That's why we always think of them as identical. If we leave BIOS and talk directly with the hardware we can differentiate between left and right shift keys with their scan code. Now this scan code is available from the keyboard data port which is 60. The keyboard generates IRQ 1 whenever a key is pressed so if we hook INT 9 and inside it read port 60 we can tell which of the shift keys was hit. Our first program will do precisely this. It will output an L if the left shift key was pressed and R if the right one was pressed. The hooking method is the same as done in the previous chapter.

Example 9.2

```

001      ; differentiate left and right shift keys with scancodes
002      [org 0x0100]
003              jmp  start
004
005      ; keyboard interrupt service routine
006      kbisr:    push ax
007                push es
008
009                mov  ax, 0xb800
010                mov  es, ax          ; point es to video memory
011
012                in   al, 0x60         ; read a char from keyboard port
013                cmp  al, 0x2a         ; is the key left shift
014                jne  nextcmp         ; no, try next comparison
015
016                mov  byte [es:0], 'L' ; yes, print L at top left
017                jmp  nomatch         ; leave interrupt routine
018
019      nextcmp:    cmp  al, 0x36         ; is the key right shift
020                jne  nomatch         ; no, leave interrupt routine
021
022                mov  byte [es:0], 'R' ; yes, print R at top left
023
024      nomatch:    mov  al, 0x20
025                out  0x20, al         ; send EOI to PIC
026
027                pop  es
028                pop  ax
029                iret
030
031      start:      xor  ax, ax
032                mov  es, ax          ; point es to IVT base
033                cli   ; disable interrupts
034                mov  word [es:9*4], kbisr ; store offset at n*4
035                mov  [es:9*4+2], cs    ; store segment at n*4+2
036                sti   ; enable interrupts

```

037			
038	11:	jmp 11	; infinite loop
033-036	CLI clears the interrupt flag to disable the interrupt system completely. The processor closes its ears and does not care about the state of the INT pin. Interrupt hooking is done in two instructions, placing the segment and placing the offset. If an interrupt comes inbetween and the vector is in an indeterminate state, the system will go to a junk address and eventually crash. So we stop all interruptions while changing a real time interrupt vector. We set the interrupt flag afterwards to reenale interrupts.		
038	The program hangs in an infinite loop. The only activity can be caused by a real time interrupt. The kbisr routine is not called from anywhere; it is only automatically invoked as a result of IRQ 1.		

When the program is executed the left and right shift keys can be distinguished with the L or R on the screen. As no action was taken for the rest of the keys, they are effectively disabled and the computer has to be rebooted. To check that the keyboard is actually disabled we change the program and add the INT 16 service 0 at the end to wait for an Esc keypress. As soon as Esc is pressed we want to terminate our program.

Example 9.3			
001	; attempt to terminate program with Esc that hooks keyboard interrupt		
002	[org 0x0100]		
003		jmp start	
004			
005-029	;;;;; COPY LINES 005-029 FROM EXAMPLE 9.2 (kbisr) ;;;;;		
030			
031	start:	xor ax, ax	
032		mov es, ax	; point es to IVT base
033		cli	; disable interrupts
034		mov word [es:9*4], kbisr	; store offset at n*4
035		mov [es:9*4+2], cs	; store segment at n*4+2
036		sti	; enable interrupts
037			
038	11:	mov ah, 0	; service 0 - get keystroke
039		int 0x16	; call BIOS keyboard service
040			
041		cmp al, 27	; is the Esc key pressed
042		jne 11	; if no, check for next key
043			
044		mov ax, 0x4c00	; terminate program
045		int 0x21	

When the program is executed the behavior is same. Esc does not work. This is because the original IRQ 1 handler was written by BIOS that read the scan code, converted into an ASCII code and stored in the keyboard buffer. The BIOS INT 16 read the key from there and gives in AL. When we hooked the keyboard interrupt BIOS is no longer in control, it has no information, it will always see the empty buffer and INT 16 will never return.

Interrupt Chaining

We can transfer control to the original BIOS ISR in the end of our routine. This way the normal functioning of INT 16 can work as well. We can retrieve the address of the BIOS routine by saving the values in vector 9 before hooking our routine. In the end of our routine we will jump to this address using a special indirect form of the JMP FAR instruction.

Example 9.4

001		; another attempt to terminate program with Esc that hooks
002		; keyboard interrupt
003		[org 0x100]
004		jmp start
005		
006	oldisr:	dd 0 ; space for saving old isr
007		
008		; keyboard interrupt service routine
009	kbisr:	push ax
010		push es
011		
012		mov ax, 0xb800
013		mov es, ax ; point es to video memory
014		
015		in al, 0x60 ; read a char from keyboard port
016		cmp al, 0x2a ; is the key left shift
017		jne nextcmp ; no, try next comparison
018		
019		mov byte [es:0], 'L' ; yes, print L at top left
020		jmp nomatch ; leave interrupt routine
021		
022	nextcmp:	cmp al, 0x36 ; is the key right shift
023		jne nomatch ; no, leave interrupt routine
024		
025		mov byte [es:0], 'R' ; yes, print R at top left
026		
027	nomatch:	; mov al, 0x20
028		; out 0x20, al
029		
030		pop es
031		pop ax
032		jmp far [cs:oldisr] ; call the original ISR
033		; iret
034		
035	start:	xor ax, ax
036		mov es, ax ; point es to IVT base
037		mov ax, [es:9*4]
038		mov [oldisr], ax ; save offset of old routine
039		mov ax, [es:9*4+2]
040		mov [oldisr+2], ax ; save segment of old routine
041		cli ; disable interrupts
042		mov word [es:9*4], kbisr ; store offset at n*4
043		mov [es:9*4+2], cs ; store segment at n*4+2
044		sti ; enable interrupts
045		
046	l1:	mov ah, 0 ; service 0 - get keystroke
047		int 0x16 ; call BIOS keyboard service
048		
049		cmp al, 27 ; is the Esc key pressed
050		jne l1 ; if no, check for next key
051		
052		mov ax, 0x4c00 ; terminate program
053		int 0x21
027-028	EOI is no longer needed as the original BIOS routine will have it at its end.	
033	IRET has been removed and an unconditional jump is introduced. At time of JMP the stack has the exact formation as was when the interrupt came. So the original BIOS routine's IRET will take control to the interrupted program. We have been careful in restoring every register we modified and retained the stack in the same form as it was at the time of entry into the routine.	

When the program is executed L and R are printed as desired and Esc terminates the program as well. Normal commands like DIR work now and shift keys still show L and R as our routine did even after the termination of our program. Now start some application like the editor, it open well but as soon as a key is pressed the computer crashes.

Actually our hookin and chaining was fine. When Esc was pressed we signaled DOS that our program has terminated. DOS will take all our

memory as a result. The routine is still in memory and functioning but the memory is free according to DOS. As soon as we load EDIT the same memory is allocated to EDIT and our routine is overwritten. Now when a key is pressed our routine's address is in the vector but at that address some new code is placed that is not intended to be an interrupt handler. That may be data or some part of the EDIT program. This results in crashing the computer.

Unhooking Interrupt

We now add the interrupt restoring part to our program. This code resets the interrupt vector to the value it had before the start of our program.

Example 9.5	
001	; terminate program with Esc that hooks keyboard interrupt
002	[org 0x100]
003	jmp start
004	
005	oldisr: dd 0 ; space for saving old isr
006	
007-032	;;;; COPY LINES 005-029 FROM EXAMPLE 9.4 (kbisr) ;;;;
033	
034	start: xor ax, ax
035	mov es, ax ; point es to IVT base
036	mov ax, [es:9*4]
037	mov [oldisr], ax ; save offset of old routine
038	mov ax, [es:9*4+2]
039	mov [oldisr+2], ax ; save segment of old routine
040	cli ; disable interrupts
041	mov word [es:9*4], kbisr ; store offset at n*4
042	mov [es:9*4+2], cs ; store segment at n*4+2
043	sti ; enable interrupts
044	
045	l1: mov ah, 0 ; service 0 - get keystroke
046	int 0x16 ; call BIOS keyboard service
047	
048	cmp al, 27 ; is the Esc key pressed
049	jne l1 ; if no, check for next key
050	
051	mov ax, [oldisr] ; read old offset in ax
052	mov bx, [oldisr+2] ; read old segment in bx
053	cli ; disable interrupts
054	mov [es:9*4], ax ; restore old offset from ax
055	mov [es:9*4+2], bx ; restore old segment from bx
056	sti ; enable interrupts
057	
058	mov ax, 0x4c00 ; terminate program
059	int 0x21

2.3. TERMINATE AND STAY RESIDENT

We change the display to show L only while the left shift is pressed and R only while the right shift is pressed to show the use of the release codes. We also changed that shift keys are not forwarded to BIOS. The effect will be visible with A and Shift-A both producing small 'a' but capslock will work.

There is one major difference from all the programs we have been writing till now. The termination is done using INT 21 service 31 instead of INT 21 service 4C. The effect is that even after termination the program is there and is legally there.

Example 9.6	
001	; TSR to show status of shift keys on top left of screen
002	[org 0x0100]
003	jmp start
004	
005	oldisr: dd 0 ; space for saving old isr
006	


```

007 ; keyboard interrupt service routine
008 kbisr: push ax
009        push es
010
011        mov ax, 0xb800
012        mov es, ax ; point es to video memory
013
014        in al, 0x60 ; read a char from keyboard port
015        cmp al, 0x2a ; has the left shift pressed
016        jne nextcmp ; no, try next comparison
017
018        mov byte [es:0], 'L' ; yes, print L at first column
019        jmp exit ; leave interrupt routine
020
021 nextcmp: cmp al, 0x36 ; has the right shift pressed
022        jne nextcmp2 ; no, try next comparison
023
024        mov byte [es:0], 'R' ; yes, print R at second column
025        jmp exit ; leave interrupt routine
026
027 nextcmp2: cmp al, 0xaa ; has the left shift released
028        jne nextcmp3 ; no, try next comparison
029
030        mov byte [es:0], ' ' ; yes, clear the first column
031        jmp exit ; leave interrupt routine
032
033 nextcmp3: cmp al, 0xb6 ; has the right shift released
034        jne nomatch ; no, chain to old ISR
035
036        mov byte [es:2], ' ' ; yes, clear the second column
037        jmp exit ; leave interrupt routine
038
039 nomatch: pop es
040        pop ax
041        jmp far [cs:oldisr] ; call the original ISR
042
043 exit:    mov al, 0x20
044        out 0x20, al ; send EOI to PIC
045
046        pop es
047        pop ax
048        iret ; return from interrupt
049
050 start:  xor ax, ax
051        mov es, ax ; point es to IVT base
052        mov ax, [es:9*4]
053        mov [oldisr], ax ; save offset of old routine
054        mov ax, [es:9*4+2]
055        mov [oldisr+2], ax ; save segment of old routine
056        cli ; disable interrupts
057        mov word [es:9*4], kbisr ; store offset at n*4
058        mov [es:9*4+2], cs ; store segment at n*4+2
059        sti ; enable interrupts
060
061        mov dx, start ; end of resident portion
062        add dx, 15 ; round up to next para
063        mov cl, 4
064        shr dx, cl ; number of paras
065        mov ax, 0x3100 ; terminate and stay resident
066        int 0x21

```

When this program is executed the command prompt immediately comes. DIR can be seen. EDIT can run and keypresses do not result in a crash. And with all that left and right shift keys shown L and R on top left of the screen while they are pressed but the shift keys do not work as usual since we did not forward the key to BIOS. This is selective chaining.

To understand Terminate and Stay Resident (TSR) programs the DOS memory formation and allocation procedure must be understood. At physical address zero is the interrupt vector table. Above it are the BIOS data area, DOS data area, IO.SYS, MSDOS.SYS and other device drivers. In the end there is COMMAND.COM command interpreter. The remaining space is called the transient program area as programs are loaded and executed in this area and the space reclaimed on their exit. A freemem pointer in DOS points where the free memory begins. When DOS loads a



program the freemem pointer is moved to the end of memory, all the available space is allocated to it, and when it exits the freemem pointer comes back to its original place thereby reclaiming all space. This action is initiated by the DOS service 4C.

The second method to legally terminate a program and give control back to DOS is using the service 31. Control is still taken back but the memory releasing part is modified. A portion of the allocated memory can be retained. So the difference in the two methods is that the freemem pointer goes back to the original place or a designated number of bytes ahead of that old position. Remember that our program crashed because the interrupt routine was overwritten. If we can tell DOS not to reclaim the memory of the interrupt routine, then it will not crash. In the last program we have told DOS to make a number of bytes resident. It becomes a part of the operation system, an extension to it. Just like DOSKEY[‡] is an extension to the operation system.

The number of paragraphs to reserve is given in the DX register. Paragraph is a unit just like byte, word, and double word. A paragraph is 16 bytes. Therefore we can reserve in multiple of 16 bytes. We write TSRs in such a way that the initialization code and data is located at the end as it is not necessary to make it resident and therefore to save space.

To calculate the number of paragraphs a label is placed after the last line that is to be made resident. The value of that label is the number of bytes needed to be made resident. A simple division by 16 will not give the correct number of paras as we want our answer to be rounded up and not down. For example 100 bytes should need 7 pages but division gives 6 and a remainder of 4. A standard technique to get rounded up integer division is to add divisor-1 to the dividend and then divide. So we add 15 to the number of bytes and then divide by 16. We use shifting for division as the divisor is a power of 2. We use a form of SHR that places the count in the CL register so that we can shift by 4 in just two instructions instead of 4 if we shift one by one.

In our program anything after start label is not needed after the program has become a TSR. We can observe that our program has become a part of DOS by giving the following command.

```
mem /c
```

[‡] DOSKEY is a TSR that shows the previous commands on the command prompt with up and down arrows and allows editing of the command

This command displays all currently loaded drivers and the current state of memory. We will be able to see our program in the list of DOS drivers.

2.4. PROGRAMMABLE INTERVAL TIMER

Another very important peripheral device is the Programmable Interval Timer (PIT), the chip numbered 8254. This chip has a precise input frequency of 1.19318 MHz. This frequency is fixed regardless of the processor clock. Inside the chip is a 16bit divisor which divides this input frequency and the output is connected to the IRQ 0 line of the PIC. The special number 0 if placed in the divisor means a divisor of 65536 and not 0. The standard divisor is 0 unless we change it. Therefore by default IRQ 0 is generated $1193180/65536=18.2$ times per second. This is called the timer tick. There is an interval of about 55ms between two timer ticks. The system time is maintained with the timer interrupt. This is the highest priority interrupt and breaks whatever is executing. Time can be maintained with this interrupt as this frequency is very precise and is part of the IBM standard.

When writing a TSR we give control back to DOS so TSR activation, reactivation and action is solely interrupt based, whether this is a hardware interrupt or a software one. Control is never given back; it must be caught, just like we caught control by hooking the keyboard interrupt. Our next example will hook the timer interrupt and display a tick count on the screen.

Example 9.7

```

001      ; display a tick count on the top right of screen
002      [org 0x0100]
003              jmp  start
004
005      tickcount:  dw  0
006
007      ; subroutine to print a number at top left of screen
008      ; takes the number to be printed as its parameter
009      printnum:   push bp
010                  mov bp, sp
011                  push es
012                  push ax
013                  push bx
014                  push cx
015                  push dx
016                  push di
017
018                  mov ax, 0xb800
019                  mov es, ax          ; point es to video base
020                  mov ax, [bp+4]      ; load number in ax
021                  mov bx, 10          ; use base 10 for division
022                  mov cx, 0           ; initialize count of digits
023
024      nextdigit:  mov dx, 0            ; zero upper half of dividend
025                  div bx              ; divide by 10
026                  add dl, 0x30        ; convert digit into ascii value
027                  push dx             ; save ascii value on stack
028                  inc cx              ; increment count of values
029                  cmp ax, 0           ; is the quotient zero
030                  jnz nextdigit       ; if no divide it again
031
032                  mov di, 140         ; point di to 70th column
033
034      nextpos:    pop dx               ; remove a digit from the stack
035                  mov dh, 0x07        ; use normal attribute
036                  mov [es:di], dx     ; print char on screen
037                  add di, 2           ; move to next screen location
038                  loop nextpos        ; repeat for all digits on stack
039
040                  pop di
041                  pop dx
042                  pop cx
043                  pop bx
044                  pop ax

```



```

045             pop  es
046             pop  bp
047             ret  2
048
049 ; timer interrupt service routine
050 timer:        push ax
051
052             inc  word [cs:tickcount]; increment tick count
053             push word [cs:tickcount]
054             call printnum           ; print tick count
055
056             mov  al, 0x20
057             out  0x20, al           ; end of interrupt
058
059             pop  ax
060             iret                    ; return from interrupt
061
062 start:        xor  ax, ax
063             mov  es, ax             ; point es to IVT base
064             cli                      ; disable interrupts
065             mov  word [es:8*4], timer; store offset at n*4
066             mov  [es:8*4+2], cs      ; store segment at n*4+2
067             sti                      ; enable interrupts
068
069             mov  dx, start           ; end of resident portion
070             add  dx, 15              ; round up to next para
071             mov  cl, 4
072             shr  dx, cl              ; number of paras
073             mov  ax, 0x3100          ; terminate and stay resident
074             int  0x21

```

When we execute the program the counter starts on the screen. Whatever we do, take directory, open EDIT, the debugger etc. the counter remains running on the screen. No one is giving control to the program; the program is getting executed as a result of timer generating INT 8 after every 55ms.

Our next example will hook both the keyboard and timer interrupts. When the shift key is pressed the tick count starts incrementing and as soon as the shift key is released the tick count stops. Both interrupt handlers are communicating through a common variable. The keyboard interrupt sets this variable while the timer interrupts modifies its behavior according to this variable.

Example 9.8

```

001 ; display a tick count while the left shift key is down
002 [org 0x0100]
003             jmp  start
004
005 seconds:     dw    0
006 timerflag:   dw    0
007 oldkb:       dd    0
008
009-049 ;;;; COPY LINES 007-047 FROM EXAMPLE 9.7 (printnum) ;;;;
050
051 ; keyboard interrupt service routine
052 kbisr:       push ax
053
054             in   al, 0x60           ; read char from keyboard port
055             cmp  al, 0x2a           ; has the left shift pressed
056             jne  nextcmp            ; no, try next comparison
057
058             cmp  word [cs:timerflag], 1; is the flag already set
059             je   exit               ; yes, leave the ISR
060
061             mov  word [cs:timerflag], 1; set flag to start printing
062             jmp  exit               ; leave the ISR
063
064 nextcmp:      cmp  al, 0xaa          ; has the left shift released
065             jne  nomatch            ; no, chain to old ISR
066
067             mov  word [cs:timerflag], 0; reset flag to stop printing

```

068		jmp exit	; leave the interrupt routine
069			
070	nomatch:	pop ax	
071		jmp far [cs:oldkb]	; call original ISR
072			
073	exit:	mov al, 0x20	
074		out 0x20, al	; send EOI to PIC
075			
076		pop ax	
077		iret	; return from interrupt
078			
079			; timer interrupt service routine
080	timer:	push ax	
081			
082		cmp word [cs:timerflag], 1	; is the printing flag set
083		jne skipall	; no, leave the ISR
084			
085		inc word [cs:seconds]	; increment tick count
086		push word [cs:seconds]	
087		call printnum	; print tick count
088			
089	skipall:	mov al, 0x20	
090		out 0x20, al	; send EOI to PIC
091			
092		pop ax	
093		iret	; return from interrupt
094			
095	start:	xor ax, ax	
096		mov es, ax	; point es to IVT base
097		mov ax, [es:9*4]	
098		mov [oldkb], ax	; save offset of old routine
099		mov ax, [es:9*4+2]	
100		mov [oldkb+2], ax	; save segment of old routine
101		cli	; disable interrupts
102		mov word [es:9*4], kbisr	; store offset at n*4
103		mov [es:9*4+2], cs	; store segment at n*4+2
104		mov word [es:8*4], timer	; store offset at n*4
105		mov [es:8*4+2], cs	; store segment at n*4+
106		sti	; enable interrupts
107			
108		mov dx, start	; end of resident portion
109		add dx, 15	; round up to next para
110		mov cl, 4	
111		shr dx, cl	; number of paras
112		mov ax, 0x3100	; terminate and stay resident
113		int 0x21	
006	This flag is one when the timer interrupt should increment and zero when it should not.		
058-059	As the keyboard controller repeatedly generates the press code if the release code does not come in a specified time, we have placed a check to not repeatedly set it to one.		
058	Another way to access TSR data is using the CS override instead of initializing DS. It is common mistake not to initialize DS and also not put in CS override in a real time interrupt handler.		

When we execute the program and the shift key is pressed, the counter starts incrementing. When the key is released the counter stops. When it is pressed again the counter resumes counting. As this is made as a TSR any other program can be loaded and will work properly alongside the TSR.

2.5. PARALLEL PORT

Computers can control external hardware through various external ports like the parallel port, the serial port, and the new additions USB and FireWire. Using this, computers can be used to control almost anything. For our examples we will use the parallel port. The parallel port has two views, the connector that the external world sees and the parallel port controller

ports through which the processor communicates with the device connected to the parallel port.

The parallel port connector is a 25pin connector called DB-25. Different pins of this connector have different meanings. Some are meaningful only with the printer[§]. This is a bidirectional port so there are some pins to take data from the processor to the parallel port and others to take data from the parallel port to the processor. Important pins for our use are the data pins from pin 2 to pin 9 that take data from the processor to the device. Pin 10, the ACK pin, is normally used by the printer to acknowledge the receipt of data and show the willingness to receive more data. Signalling this pin generates IRQ 7 if enabled in the PIC and in the parallel port controller. Pin 18-25 are ground and must be connected to the external circuit ground to provide the common reference point otherwise they won't understand each other voltage levels. Like the datum point in a graph this is the datum point of an electrical circuit. The remaining pins are not of our concern in these examples.

This is the external view of the parallel port. The processor cannot see these pins. The processor uses the I/O ports of the parallel port controller. The first parallel port LPT1** has ports designated from 378 to 37A. The first port 378 is the data port. If we use the OUT instruction on this port, 1 bits result in a 5V signal on the corresponding pin and a 0 bits result in a 0V signal on the corresponding pin.

Port 37A is the control port. Our interest is with bit 4 of this port which enables the IRQ 7 triggering by the ACK pin. We have attached a circuit that connects 8 LEDs with the parallel port pins. The following examples sends the scancode of the key pressed to the parallel port so that it is visible on LEDs.

Example 9.9	
001	; show scancode on external LEDs connected through parallel port
002	[org 0x0100]
003	jmp start
004	
005	oldisr: dd 0 ; space for saving old ISR
006	
007	; keyboard interrupt service routine
008	kbisr: push ax
009	push dx
010	
011	in al, 0x60 ; read char from keyboard port
012	mov dx, 0x378
013	out dx, al ; write char to parallel port
014	
015	pop ax
016	pop dx
017	jmp far [cs:oldisr] ; call original ISR
018	
019	start: xor ax, ax
020	mov es, ax ; point es to IVT base
021	mov ax, [es:9*4]
022	mov [oldisr], ax ; save offset of old routine
023	mov ax, [es:9*4+2]
024	mov [oldisr+2], ax ; save segment of old routine
025	cli ; disable interrupts
026	mov word [es:9*4], kbisr ; store offset at n*4
027	mov [es:9*4+2], cs ; store segment at n*4+2
028	sti ; enable interrupts
029	
030	mov dx, start ; end of resident portion
031	add dx, 15 ; round up to next para

[§] The parallel port is most commonly used with the printer. However some new printers have started using the USB port.

** Older computer had more than one parallel port named LPT2 and having ports from 278-27A.

```

032      mov     cl, 4
033      shr     dx, cl           ; number of paras
034      mov     ax, 0x3100      ; terminate and stay resident
035      int     0x21

```

The following example uses the same LED circuit and sends data such that LEDs switch on and off turn by turn so that it looks like light is moving back and forth.

Example 9.10	
001	; show lights moving back and forth on external LEDs
002	[org 0x0100]
003	jmp start
004	
005	signal: db 1 ; current state of lights
006	direction: db 0 ; current direction of motion
007	
008	; timer interrupt service routine
009	timer: push ax
010	push dx
011	push ds
012	
013	push cs
014	pop ds ; initialize ds to data segment
015	
016	cmp byte [direction], 1; are moving in right direction
017	je moveright ; yes, go to shift right code
018	
019	shl byte [signal], 1 ; shift left state of lights
020	jnc output ; no jump to change direction
021	
022	mov byte [direction], 1; change direction to right
023	mov byte [signal], 0x80; turn on left most light
024	jmp output ; proceed to send signal
025	
026	moveright: shr byte [signal], 1 ; shift right state of lights
027	jnc output ; no jump to change direction
028	
029	mov byte [direction], 0; change direction to left
030	mov byte [signal], 1 ; turn on right most light
031	
032	output: mov al, [signal] ; load lights state in al
033	mov dx, 0x378 ; parallel port data port
034	out dx, al ; send light state of port
035	
036	mov al, 0x20
037	out 0x20, al ; send EOI on PIC
038	
039	pop ds
040	pop dx
041	pop ax
042	iret ; return from interrupt
043	
044	start: xor ax, ax
045	mov es, ax ; point es to IVT base
046	cli ; disable interrupts
047	mov word [es:8*4], timer ; store offset at n*4
048	mov [es:8*4+2], cs ; store segment at n*4+2
049	sti ; enable interrupts
050	
051	mov dx, start ; end of resident portion
052	add dx, 15 ; round up to next para
053	mov cl, 4
054	shr dx, cl ; number of paras
055	mov ax, 0x3100 ; terminate and stay resident
056	int 0x21

We will now use the parallel port to control a slightly complicated circuit. This time we will also use the parallel port interrupt. We are using a 220 V bulb with AC input. AC current is 50Hz sine wave. We have made our circuit such that it triggers the parallel port interrupt whenever the sine wave

crosses zero. We have control of passing the AC current to the bulb. We control it such that in every cycle only a fixed percentage of time the current passes on to the bulb. Using this we can control the intensity or glow of the bulb.

Our first example will slowly turn on the bulb by increasing the power provided using the mechanism just described.

Example 9.11

```

001 ; slowly turn on a bulb by gradually increasing the power provided
002 [org 0x0100]
003 jmp start
004
005 flag: db 0 ; next time turn on or turn off
006 stop: db 0 ; flag to terminate the program
007 divider: dw 11000 ; divider for minimum intensity
008 oldtimer: dd 0 ; space for saving old isr
009
010 ; timer interrupt service routine
011 timer: push ax
012 push dx
013
014 cmp byte [cs:flag], 0 ; are we here to turn off
015 je switchoff ; yes, go to turn off code
016
017 switchon: mov al, 1
018 mov dx, 0x378
019 out dx, al ; no, turn the bulb on
020
021 mov ax, 0x0100
022 out 0x40, al ; set timer divisor LSB to 0
023 mov al, ah
024 out 0x40, al ; set timer divisor MSB to 1
025 mov byte [cs:flag], 0 ; flag next timer to switch off
026 jmp exit ; leave the interrupt routine
027
028 switchoff: xor ax, ax
029 mov dx, 0x378
030 out dx, al ; turn the bulb off
031
032 exit: mov al, 0x20
033 out 0x20, al ; send EOI to PIC
034
035 pop dx
036 pop ax
037 iret ; return from interrupt
038
039 ; parallel port interrupt service routine
040 parallel: push ax
041
042 mov al, 0x30 ; set timer to one shot mode
043 out 0x43, al
044
045 cmp word [cs:divider], 100; is the current divisor 100
046 je stopit ; yes, stop
047
048 sub word [cs:divider], 10; decrease the divisor by 10
049 mov ax, [cs:divider]
050 out 0x40, al ; load divisor LSB in timer
051 mov al, ah
052 out 0x40, al ; load divisor MSB in timer
053 mov byte [cs:flag], 1 ; flag next timer to switch on
054
055 mov al, 0x20
056 out 0x20, al ; send EOI to PIC
057 pop ax
058 iret ; return from interrupt
059
060 stopit: mov byte [stop], 1 ; flag to terminate the program
061 mov al, 0x20
062 out 0x20, al ; send EOI to PIC
063 pop ax
064 iret ; return from interrupt
065

```

```

066      start:      xor  ax, ax
067                  mov  es, ax          ; point es to IVT base
068                  mov  ax, [es:0x08*4]
069                  mov  [oldtimer], ax  ; save offset of old routine
070                  mov  ax, [es:0x08*4+2]
071                  mov  [oldtimer+2], ax ; save segment of old routine
072                  cli                      ; disable interrupts
073                  mov  word [es:0x08*4], timer ; store offset at n*4
074                  mov  [es:0x08*4+2], cs  ; store segment at n*4+2
075                  mov  word [es:0x0F*4], parallel ; store offset at n*4
076                  mov  [es:0x0F*4+2], cs  ; store segment at n*4+2
077                  sti                      ; enable interrupts
078
079                  mov  dx, 0x37A
080                  in   al, dx             ; parallel port control register
081                  or   al, 0x10           ; turn interrupt enable bit on
082                  out  dx, al             ; write back register
083
084                  in   al, 0x21           ; read interrupt mask register
085                  and  al, 0x7F          ; enable IRQ7 for parallel port
086                  out  0x21, al          ; write back register
087
088      recheck:     cmp  byte [stop], 1    ; is the termination flag set
089                  jne  recheck           ; no, check again
090
091                  mov  dx, 0x37A
092                  in   al, dx             ; parallel port control register
093                  and  al, 0xEF           ; turn interrupt enable bit off
094                  out  dx, al             ; write back register
095
096                  in   al, 0x21           ; read interrupt mask register
097                  or   al, 0x80           ; disable IRQ7 for parallel port
098                  out  0x21, al          ; write back register
099
100                  cli                      ; disable interrupts
101                  mov  ax, [oldtimer]     ; read old timer ISR offset
102                  mov  [es:0x08*4], ax    ; restore old timer ISR offset
103                  mov  ax, [oldtimer+2]   ; read old timer ISR segment
104                  mov  [es:0x08*4+2], ax  ; restore old timer ISR segment
105                  sti                      ; enable interrupts
106
107                  mov  ax, 0x4c00         ; terminate program
108                  int  0x21

```

The next example is simply the opposite of the previous. It slowly turns the bulb off from maximum glow to no glow.

Example 9.12

```

001      ; slowly turn off a bulb by gradually decreasing the power provided
002      [org 0x0100]
003                  jmp  start
004
005      flag:        db   0                ; next time turn on or turn off
006      stop:        db   0                ; flag to terminate the program
007      divider:     dw   0                ; divider for maximum intensity
008      oldtimer:    dd   0                ; space for saving old isr
009
010-037  ;;;; COPY LINES 009-036 FROM EXAMPLE 9.11 (timer) ;;;;
038
039      ; parallel port interrupt service routine
040      parallel:     push ax
041
042                  mov  al, 0x30           ; set timer to one shot mode
043                  out  0x43, al
044
045                  cmp  word [cs:divider], 11000; current divisor is 11000
046                  je   stopit            ; yes, stop
047
048                  add  word [cs:divider], 10; increase the divisor by 10
049                  mov  ax, [cs:divider]
050                  out  0x40, al           ; load divisor LSB in timer
051                  mov  al, ah
052                  out  0x40, al           ; load divisor MSB in timer

```

```

053             mov     byte [cs:flag], 1    ; flag next timer to switch on
054
055             mov     al, 0x20
056             out     0x20, al              ; send EOI to PIC
057             pop     ax
058             iret                     ; return from interrupt
059
060 stopit:      mov     byte [stop], 1       ; flag to terminate the program
061             mov     al, 0x20
062             out     0x20, al              ; send EOI to PIC
063             pop     ax
064             iret                     ; return from interrupt
065
066 start:      xor     ax, ax
067             mov     es, ax                ; point es to IVT base
068             mov     ax, [es:0x08*4]
069             mov     [oldtimer], ax        ; save offset of old routine
070             mov     ax, [es:0x08*4+2]
071             mov     [oldtimer+2], ax      ; save segment of old routine
072             cli                     ; disable interrupts
073             mov     word [es:0x08*4], timer ; store offset at n*4
074             mov     [es:0x08*4+2], cs    ; store segment at n*4+2
075             mov     word [es:0x0F*4], parallel ; store offset at n*4
076             mov     [es:0x0F*4+2], cs    ; store segment at n*4+2
077             sti                     ; enable interrupts
078
079             mov     dx, 0x37A
080             in      al, dx                ; parallel port control register
081             or      al, 0x10              ; turn interrupt enable bit on
082             out     dx, al                ; write back register
083
084             in      al, 0x21              ; read interrupt mask register
085             and     al, 0x7F              ; enable IRQ7 for parallel port
086             out     0x21, al              ; write back register
087
088 recheck:    cmp     byte [stop], 1       ; is the termination flag set
089             jne     recheck               ; no, check again
090
091             mov     dx, 0x37A
092             in      al, dx                ; parallel port control register
093             and     al, 0xEF              ; turn interrupt enable bit off
094             out     dx, al                ; write back register
095
096             in      al, 0x21              ; read interrupt mask register
097             or      al, 0x80              ; disable IRQ7 for parallel port
098             out     0x21, al              ; write back register
099
100            cli                     ; disable interrupts
101            mov     ax, [oldtimer]         ; read old timer ISR offset
102            mov     [es:0x08*4], ax        ; restore old timer ISR offset
103            mov     ax, [oldtimer+2]      ; read old timer ISR segment
104            mov     [es:0x08*4+2], ax     ; restore old timer ISR segment
105            sti                     ; enable interrupts
106
107            mov     ax, 0x4c00              ; terminate program
108            int     0x21

```

This example is a mix of the previous two. Here we can increase the bulb intensity with F11 and decrease it with F12.

Example 9.13

```

001             ; control external bulb intensity with F11 and F12
002             [org 0x0100]
003             jmp     start
004
005 flag:        db      0                    ; next time turn on or turn off
006 divider:    dw      100                  ; initial timer divider
007 oldkb:      dd      0                    ; space for saving old ISR
008
009-036        ;;;; COPY LINES 009-036 FROM EXAMPLE 9.11 (timer) ;;;;
037
038             ; keyboard interrupt service routine
039 kbisr:      push    ax

```

```

040
041         in    al, 0x60
042         cmp   al, 0x57
043         jne   nextcmp
044         cmp   word [cs:divider], 11000
045         je    exitkb
046         add   word [cs:divider], 100
047         jmp   exitkb
048
049     nextcmp:    cmp   al, 0x58
050                 jne   chain
051                 cmp   word [cs:divider], 100
052                 je    exitkb
053                 sub   word [cs:divider], 100
054                 jmp   exitkb
055
056     exitkb:     mov   al, 0x20
057                 out   0x20, al
058
059                 pop   ax
060                 iret
061
062     chain:      pop   ax
063                 jmp   far [cs:oldkb]
064
065     ; parallel port interrupt service routine
066     parallel:   push  ax
067
068                 mov   al, 0x30           ; set timer to one shot mode
069                 out   0x43, al
070
071                 mov   ax, [cs:divider]
072                 out   0x40, al           ; load divisor LSB in timer
073                 mov   al, ah
074                 out   0x40, al           ; load divisor MSB in timer
075                 mov   byte [cs:flag], 1 ; flag next timer to switch on
076
077                 mov   al, 0x20
078                 out   0x20, al           ; send EOI to PIC
079                 pop   ax
080                 iret                     ; return from interrupt
081
082     start:      xor   ax, ax
083                 mov   es, ax             ; point es to IVT base
084                 mov   ax, [es:0x09*4]
085                 mov   [oldkb], ax        ; save offset of old routine
086                 mov   ax, [es:0x09*4+2]
087                 mov   [oldkb+2], ax      ; save segment of old routine
088                 cli                     ; disable interrupts
089                 mov   word [es:0x08*4], timer ; store offset at n*4
090                 mov   [es:0x08*4+2], cs  ; store segment at n*4+2
091                 mov   word [es:0x09*4], kbisr ; store offset at n*4
092                 mov   [es:0x09*4+2], cs  ; store segment at n*4+2
093                 mov   word [es:0x0F*4], parallel ; store offset at n*4
094                 mov   [es:0x0F*4+2], cs  ; store segment at n*4+2
095                 sti                     ; enable interrupts
096
097                 mov   dx, 0x37A
098                 in    al, dx             ; parallel port control register
099                 or    al, 0x10           ; turn interrupt enable bit on
100                 out   dx, al            ; write back register
101
102                 in    al, 0x21           ; read interrupt mask register
103                 and   al, 0x7F           ; enable IRQ7 for parallel port
104                 out   0x21, al          ; write back register
105
106                 mov   dx, start          ; end of resident portion
107                 add   dx, 15             ; round up to next para
108                 mov   cl, 4
109                 shr   dx, cl             ; number of paras
110                 mov   ax, 0x3100        ; terminate and stay resident
111                 int   0x21

```

EXERCISES

1. Suggest a reason for the following. The statements are all true.
 - a. We should disable interrupts while hooking interrupt 8h. I.e. while placing its segment and offset in the interrupt vector table.
 - b. We need not do this for interrupt 80h.
 - c. We need not do this when hooking interrupt 8h from inside the interrupt handler of interrupt 80h.
 - d. We should disable interrupts while we are changing the stack (SS and SP).
 - e. EOI is not sent from an interrupt handler which does interrupt chaining.
 - f. If no EOI is sent from interrupt 9h and no chaining is done, interrupt 8h still comes if the interrupt flag is on.
 - g. After getting the size in bytes by putting a label at the end of a COM TSR, 0fh is added before dividing by 10h.
 - h. Interrupts are disabled but divide by zero interrupt still comes.
2. If no hardware interrupts are coming, what are all possible reasons?
3. Write a program to make an asterisks travel the border of the screen, from upper left to upper right to lower right to lower left and back to upper left indefinitely, making each movement after one second.
4. [Musical Arrow] Write a TSR to make an arrow travel the border of the screen from top left to top right to bottom right to bottom left and back to top left at the speed of 36.4 locations per second. The arrow should not destroy the data beneath it and should be restored as soon as the arrow moves forward.
 The arrow head should point in the direction of movement using the characters > V < and ^. The journey should be accompanied by a different tone from the pc speaker for each side of the screen. Do interrupt chaining so that running the TSR 10 times produces 10 arrows travelling at different locations.
 HINT: At the start you will need to reprogram channel 0 for 36.4 interrupts per second, double the normal. You will have to reprogram channel 2 at every direction change, though you can enable the speaker once at the very start.
5. In the above TSR hook the keyboard interrupt as well and check if 'q' is pressed. If not chain to the old interrupt, if yes restore everything and remove the TSR from memory. The effect should be that pressing 'q' removes one moving arrow. If you do interrupt chaining when pressing 'q' as well, it will remove all arrows at once.
6. Write a TSR to rotate the screen (scroll up and copy the old top most line to the bottom) while F10 is pressed. The screen will keep rotating while F10 is pressed at 18.2 rows per second. As soon as F10 is released the rotation should stop and the original screen restored. A secondary buffer of only 160 bytes (one line of screen) can be used.
7. Write a TSR that hooks software interrupt 0x80 and the timer interrupt. The software interrupt is called by other programs with the address of a far function in ES:DI and the number of timer ticks after which to call back that function in CX. The interrupt records this information and returns to the caller. The function will actually be called by the timer interrupt after the desired number of ticks. The maximum number of functions and their ticks can be fixed to 8.
8. Write a TSR to clear the screen when CTRL key is pressed and restore it when it is released.
9. Write a TSR to disable all writes to the hard disk when F10 is pressed and re-enable when pressed again like a toggle.

HINT: To write to the hard disk programs call the BIOS service INT 0x13 with AH=3.

10. Write a keyboard interrupt handler that disables the timer interrupt (no timer interrupt should come) while Q is pressed. It should be re-enabled as soon as Q is released.
11. Write a TSR to calculate the current typing speed of the user. Current typing speed is the number of characters typed by the user in the last five seconds. The speed should be represented by printing asterisks at the right border (80th column) of the screen starting from the upper right to the lower right corner (growing downwards). Draw n asterisks if the user typed n characters in the last five seconds. The count should be updated every second.
12. Write a TSR to show a clock in the upper right corner of the screen in the format HH:MM:SS.DD where HH is hours in 24 hour format, MM is minutes, SS is seconds and DD is hundredth of second. The clock should beep twice for one second each time with half a second interval in between at the start of every minute at a frequency of your choice.

HINT: IBM PC uses a Real Time Clock (RTC) chip to keep track of time while switched off. It provides clock and calendar functions through its two I/O ports 70h and 71h. It is used as follows:

```
        mov  al, <command>
        out  0x70, al          ; command byte written at first port
        jmp  D1                ; waste one instruction time
D1:      in   al, 0x71          ; result of command is in AL now
```

Following are few commands

- 00 Get current second
- 02 Get current minute
- 04 Get current hour

All numbers returned by RTC are in BCD. E.g. if it is 6:30 the second and third command will return 0x30 and 0x06 respectively in al.

3

Debug Interrupts

3.1. DEBUGGER USING SINGLE STEP INTERRUPT

The use of the trap flag has been deferred till now. The three flags not used for mathematical operations are the direction flag, the interrupt flag and the trap flag. The direction and interrupt flags have been previously discussed.

If the interrupt flag is set, the after every instruction a type 1 interrupt will be automatically generated. When the IVT and reserved interrupts were discussed this was named as the single step interrupt. This is like the divide by zero interrupt which was never explicitly invoked but it came itself. The single step interrupt behaves in the same manner.

The debugger is made using this interrupt. It allows one instruction to be executed and then return control to us. It has its display code and its code to wait for the key in the INT 1 handler. Therefore after every instruction the values of all registers are shown and the debugger waits for a key. Another interrupt used by the debugger is the breakpoint interrupt INT 3. Apart from single stepping debugger has the breakpoint feature. INT 3 is used for this feature. INT 3 has a single byte opcode so it can replace any instruction. To put a breakpoint the instruction is replaced with INT 3 opcode and restored in the INT 3 handler. The INT 3 opcode is again placed by a single step interrupt that is set up for this purpose after the replaced instruction has been executed.

There is no instruction to set or clear the trap flag like there are instructions for the interrupt and direction flags. We use two special instructions PUSHF and POPF to push and pop the flag from the stack. We use PUSHF to place flags on the stack, change TF in this image on the stack and then reload into the flags register with POPF. The single step interrupt will come after the first instruction after POPF. The interrupt mechanism automatically clears IF and TF otherwise there would be an infinite recursion of the single step interrupt. The TF is set in the flags on the stack so another interrupt will come after one more instruction is executed after the return of the interrupt.

The following example is a very elementary debugger using the trap flag and the single step interrupt.

Example 10.1

```
001 ; single stepping using the trap flag and single step interrupt
002 [org 0x0100]
003         jmp start
004
005 flag:      db    0                ; flag whether a key pressed
006 oldisr:    dd    0                ; space for saving old ISR
007 names:     db    'FL =CS =IP =BP =AX =BX =CX =DX =SI =DI =DS =ES ='
008
009-026 ;;;; COPY LINES 008-025 FROM EXAMPLE 6.2 (clrscr) ;;;;
027
028 ; subroutine to print a number on screen
029 ; takes the row no, column no, and number to be printed as parameters
030 printnum:  push bp
031            mov  bp, sp
032            push es
033            push ax
034            push bx
035            push cx
```

```

036          push dx
037          push di
038
039          mov di, 80          ; load di with columns per row
040          mov ax, [bp+8]      ; load ax with row number
041          mul di              ; multiply with columns per row
042          mov di, ax          ; save result in di
043          add di, [bp+6]      ; add column number
044          shl di, 1           ; turn into byte count
045          add di, 8           ; to end of number location
046
047          mov ax, 0xb800
048          mov es, ax          ; point es to video base
049          mov ax, [bp+4]      ; load number in ax
050          mov bx, 16          ; use base 16 for division
051          mov cx, 0           ; initialize count of digits
052
053  nextdigit:  mov dx, 0        ; zero upper half of dividend
054              div bx          ; divide by 10
055              add dl, 0x30     ; convert digit into ascii value
056              cmp dl, 0x39    ; is the digit an alphabet
057              jbe skipalpha   ; no, skip addition
058              add dl, 7        ; yes, make in alphabet code
059  skipalpha:  mov dh, 0x07     ; attach normal attribute
060              mov [es:di], dx  ; print char on screen
061              sub di, 2        ; to previous screen location
062              loop nextdigit   ; if no divide it again
063
064          pop di
065          pop dx
066          pop cx
067          pop bx
068          pop ax
069          pop es
070          pop bp
071          ret 6
072
073          ; subroutine to print a string
074          ; takes row no, column no, address of string, and its length
075          ; as parameters
076  printstr:  push bp
077              mov bp, sp
078              push es
079              push ax
080              push bx
081              push cx
082              push dx
083              push si
084              push di
085
086              mov ax, 0xb800
087              mov es, ax      ; point es to video base
088
089              mov di, 80      ; load di with columns per row
090              mov ax, [bp+8]   ; load ax with row number
091              mul di           ; multiply with columns per row
092              mov di, ax       ; save result in di
093              add di, [bp+6]   ; add column number
094              shl di, 1        ; turn into byte count
095
096              mov si, [bp+6]   ; string to be printed
097              mov cx, [bp+4]   ; length of string
098              mov ah, 0x07     ; normal attribute is fixed
099
100  nextchar:  mov al, [si]      ; load next char of string
101              mov [es:di], ax  ; show next char on screen
102              add di, 2        ; move to next screen location
103              add si, 1         ; move to next char
104              loop nextchar    ; repeat the operation cx times
105
106          pop di
107          pop si
108          pop dx
109          pop cx
110          pop bx
111          pop ax

```

```

112             pop  es
113             pop  bp
114             ret   8
115
116 ; keyboard interrupt service routine
117 kbisr:      push ax
118
119             in    al, 0x60          ; read a char from keyboard port
120             test  al, 0x80          ; is it a press code
121             jnz   skipflag         ; no, leave the interrupt
122             add   byte [cs:flag], al ; yes, set flag to proceed
123
124 skipflag:    mov   al, 0x20
125             out   0x20, al
126             pop   ax
127             iret
128
129 ; single step interrupt service routine
130 trapisr:    push bp
131             mov   bp, sp           ; to read cs, ip and flags
132             push  ax
133             push  bx
134             push  cx
135             push  dx
136             push  si
137             push  di
138             push  ds
139             push  es
140
141             sti                     ; waiting for keyboard interrupt
142             push  cs
143             pop   ds               ; initialize ds to data segment
144
145             mov   byte [flag], 0   ; set flag to wait for key
146             call  clrscr           ; clear the screen
147
148             mov   si, 6            ; first register is at bp+6
149             mov   cx, 12           ; total 12 registers to print
150             mov   ax, 0            ; start from row 0
151             mov   bx, 5            ; print at column 5
152
153 13:         push  ax               ; row number
154             push  bx               ; column number
155             mov   dx, [bp+si]
156             push  dx               ; number to be printed
157             call  printnum         ; print the number
158             sub   si, 2            ; point to next register
159             inc   ax               ; next row number
160             loop  13              ; repeat for the 12 registers
161
162             mov   ax, 0            ; start from row 0
163             mov   bx, 0            ; start from column 0
164             mov   cx, 12           ; total 12 register names
165             mov   si, 4            ; each name length is 4 chars
166             mov   dx, names        ; offset of first name in dx
167
168 11:         push  ax               ; row number
169             push  bx               ; column number
170             push  dx               ; offset of string
171             push  si               ; length of string
172             call  printstr         ; print the string
173             add   dx, 4            ; point to start of next string
174             inc   ax               ; new row number
175             loop  11              ; repeat for 12 register names
176
177 keywait:    cmp   byte [flag], 0   ; has a key been pressed
178             je    keywait         ; no, check again
179
180             pop   es
181             pop   ds
182             pop   di
183             pop   si
184             pop   dx
185             pop   cx
186             pop   bx
187             pop   ax
188

```

```

189         pop bp
190         iret
191
192     start:    xor ax, ax
193             mov es, ax                ; point es to IVT base
194             mov ax, [es:9*4]
195             mov [oldisr], ax          ; save offset of old routine
196             mov ax, [es:9*4+2]
197             mov [oldisr+2], ax        ; save segment of old routine
198             mov word [es:1*4], trapisr ; store offset at n*4
199             mov [es:1*4+2], cs        ; store segment at n*4+2
200             cli                       ; disable interrupts
201             mov word [es:9*4], kbisr  ; store offset at n*4
202             mov [es:9*4+2], cs        ; store segment at n*4+2
203             sti                       ; enable interrupts
204
205             pushf                     ; save flags on stack
206             pop ax                   ; copy flags in ax
207             or ax, 0x100              ; set bit corresponding to TF
208             push ax                   ; save ax on stack
209             popf                     ; reload into flags register
210
211     ; the trap flag bit is on now, INT 1 will come after next instruction
212     ; sample code to check the working of our elementary debugger
213     mov ax, 0
214     mov bx, 0x10
215     mov cx, 0x20
216     mov dx, 0x40
217
218     12:      inc ax
219             add bx, 2
220             dec cx
221             sub dx, 2
222             jmp 12

```

3.2. DEBUGGER USING BREAKPOINT INTERRUPT

We now write a debugger using INT 3. This debugger stops at the same point everytime where the breakpoint has been set up unlike the previous one which stopped at every instruction. The single step interrupt in this example is used only to restore the breakpoint interrupt which was removed by the breakpoint interrupt handler temporarily so that the original instruction can be executed.

Example 10.2

```

001     ; elementary debugger using breakpoint interrupt
002     [org 0x0100]
003             jmp start
004
005     flag:      db 0                    ; flag whether a key pressed
006     oldisr:    dd 0                    ; space for saving old ISR
007     names:     db 'FL =CS =IP =BP =AX =BX =CX =DX =SI =DI =DS =ES ='
008     opcode:     db 0
009     opcodepos: dw 0
010
011-028    ;;;; COPY LINES 008-025 FROM EXAMPLE 6.2 (clrscr) ;;;;
029-072    ;;;; COPY LINES 028-071 FROM EXAMPLE 10.1 (prntnum) ;;;;
073-114    ;;;; COPY LINES 073-114 FROM EXAMPLE 10.1 (prntstr) ;;;;
115-127    ;;;; COPY LINES 116-128 FROM EXAMPLE 10.1 (kbisr) ;;;;
128
129     ; single step interrupt service routine
130     trapisr:   push bp
131             mov bp, sp
132             push ax
133             push di
134             push ds
135             push es
136
137             push cs
138             pop ds                    ; initialize ds to data segment

```

```

139
140     mov ax, [bp+4]
141     mov es, ax           ; load interrupted segment in es
142     mov di, [opcodepos] ; load saved opcode position
143     mov byte [es:di], 0xCC ; reset the opcode to INT3
144     and word [bp+6], 0xFEFF ; clear TF in flags on stack
145
146     pop es
147     pop ds
148     pop di
149     pop ax
150     pop bp
151     iret
152
153 ; breakpoint interrupt service routine
154 debugisr: push bp
155           mov bp, sp           ; to read cs, ip and flags
156           push ax
157           push bx
158           push cx
159           push dx
160           push si
161           push di
162           push ds
163           push es
164
165           sti                 ; waiting for keyboard interrupt
166           push cs
167           pop ds             ; initialize ds to data segment
168
169           mov ax, [bp+4]
170           mov es, ax         ; load interrupted segment in es
171           dec word [bp+2]    ; decrement the return address
172           mov di, [bp+2]    ; read the return address in di
173           mov word [opcodepos], di ; remember the return position
174           mov al, [opcode]  ; load the original opcode
175           mov [es:di], al   ; restore original opcode there
176
177           mov byte [flag], 0 ; set flag to wait for key
178           call clrscr       ; clear the screen
179
180           mov si, 6         ; first register is at bp+6
181           mov cx, 12        ; total 12 registers to print
182           mov ax, 0         ; start from row 0
183           mov bx, 5         ; print at column 5
184
185 13:      push ax             ; row number
186           push bx           ; column number
187           mov dx, [bp+si]
188           push dx           ; number to be printed
189           call printnum     ; print the number
190           sub si, 2         ; point to next register
191           inc ax            ; next row number
192           loop 13          ; repeat for the 12 registers
193
194           mov ax, 0         ; start from row 0
195           mov bx, 0         ; start from column 0
196           mov cx, 12        ; total 12 register names
197           mov si, 4         ; each name length is 4 chars
198           mov dx, names     ; offset of first name in dx
199
200 11:      push ax             ; row number
201           push bx           ; column number
202           push dx           ; offset of string
203           push si           ; length of string
204           call printstr     ; print the string
205           add dx, 4         ; point to start of next string
206           inc ax            ; new row number
207           loop 11          ; repeat for 12 register names
208
209           or word [bp+6], 0x0100 ; set TF in flags image on stack
210
211 keywait: cmp byte [flag], 0 ; has a key been pressed
212           je keywait       ; no, check again
213
214           pop es

```

```

215         pop ds
216         pop di
217         pop si
218         pop dx
219         pop cx
220         pop bx
221         pop ax
222         pop bp
223         iret
224
225     start:    xor     ax, ax
226             mov     es, ax                ; point es to IVT base
227             mov     word [es:1*4], trapisr ; store offset at n*4
228             mov     [es:1*4+2], cs        ; store segment at n*4+2
229             mov     word [es:3*4], debugisr ; store offset at n*4
230             mov     [es:3*4+2], cs        ; store segment at n*4+2
231             cli                     ; disable interrupts
232             mov     word [es:9*4], kbisr  ; store offset at n*4
233             mov     [es:9*4+2], cs        ; store segment at n*4+2
234             sti                     ; enable interrupts
235
236             mov     si, 12                ; load breakpoint position in si
237             mov     al, [cs:si]           ; read opcode at that position
238             mov     [opcode], al         ; save opcode for later use
239             mov     byte [cs:si], 0xCC    ; change opcode to INT3
240
241             ; breakpoint is set now, INT3 will come at 12 on every iteration
242             ; sample code to check the working of our elementary debugger
243             mov     ax, 0
244             mov     bx, 0x10
245             mov     cx, 0x20
246             mov     dx, 0x40
247
248     12:      inc     ax
249             add     bx, 2
250             dec     cx
251             sub     dx, 2
252             jmp     12

```


4

Multitasking

4.1. CONCEPTS OF MULTITASKING

To experience the power of assembly language we introduce how to implement multitasking. We observed in the debugger that our thread of instructions was broken by the debugger; it got the control, used all registers, displayed an elaborate interface, waited for the key, and then restored processor state to what was immediately before interruption. Our program resumed as if nothing happened. The program execution was in the same logical flow.

If we have two different programs A and B. Program A is broken, its state saved, and returned to B instead of A. By looking at the instruction set, we can immediately say that nothing can stop us from doing that. IRET will return to whatever CS and IP it finds on the stack. Now B is interrupted somehow, its state saved, and we return back to A. A will have no way of knowing that it was interrupted as its entire environment has been restored. It never knew the debugger took control when it was debugged. It still has no way of gaining this knowledge. If this work of breaking and restoring programs is done at high speed the user will feel that all the programs are running at the same time where actually they are being switched to and forth at high speed.

In essence multitasking is simple, even though we have to be extremely careful when implementing it. The environment of a program in the very simple case is all its registers and stack. We will deal with stack later. Now to get control from the program without the program knowing about it, we can use the IRQ 0 highest priority interrupt that is periodically coming to the processor.

Now we present a very basic example of multitasking. We have two subroutines written in assembly language. All the techniques discussed here are applicable to code written in higher level languages as well. However the code to control this multitasking cannot be easily written in a higher level language so we write it in assembly language. The two subroutines rotate bars by changing characters at the two corners of the screen and have infinite loops. By hooking the timer interrupt and saving and restoring the registers of the tasks one by one, it appears that both tasks are running simultaneously.

Example 11.1

```
001 ; elementary multitasking of two threads
002 [org 0x0100]
003     jmp start
004
005 ; ax,bx,ip,cs,flags storage area
006 taskstates: dw 0, 0, 0, 0, 0 ; task0 regs
007             dw 0, 0, 0, 0, 0 ; task1 regs
008             dw 0, 0, 0, 0, 0 ; task2 regs
009
010 current:    db 0 ; index of current task
011 chars:     db '\|/-' ; shapes to form a bar
012
013 ; one task to be multitasked
014 taskone:   mov al, [chars+bx] ; read the next shape
015           mov [es:0], al ; write at top left of screen
```

```

016             inc bx             ; increment to next shape
017             and bx, 3           ; taking modulus by 4
018             jmp taskone        ; infinite task
019
020 ; second task to be multitasked
021 tasktwo:     mov al, [chars+bx] ; read the next shape
022             mov [es:158], al    ; write at top right of screen
023             inc bx             ; increment to next shape
024             and bx, 3           ; taking modulus by 4
025             jmp tasktwo        ; infinite task
026
027 ; timer interrupt service routine
028 timer:      push ax
029             push bx
030
031             mov bl, [cs:current] ; read index of current task
032             mov ax, 10           ; space used by one task
033             mul bl               ; multiply to get start of task
034             mov bx, ax           ; load start of task in bx
035
036             pop ax               ; read original value of bx
037             mov [cs:taskstates+bx+2], ax ; space for current task
038             pop ax               ; read original value of ax
039             mov [cs:taskstates+bx+0], ax ; space for current task
040             pop ax               ; read original value of ip
041             mov [cs:taskstates+bx+4], ax ; space for current task
042             pop ax               ; read original value of cs
043             mov [cs:taskstates+bx+6], ax ; space for current task
044             pop ax               ; read original value of flags
045             mov [cs:taskstates+bx+8], ax ; space for current task
046
047             inc byte [cs:current] ; update current task index
048             cmp byte [cs:current], 3 ; is task index out of range
049             jne skipreset        ; no, proceed
050             mov byte [cs:current], 0 ; yes, reset to task 0
051
052 skipreset:   mov bl, [cs:current] ; read index of current task
053             mov ax, 10           ; space used by one task
054             mul bl               ; multiply to get start of task
055             mov bx, ax           ; load start of task in bx
056
057             mov al, 0x20
058             out 0x20, al         ; send EOI to PIC
059
060             push word [cs:taskstates+bx+8] ; flags of new task
061             push word [cs:taskstates+bx+6] ; cs of new task
062             push word [cs:taskstates+bx+4] ; ip of new task
063             mov ax, [cs:taskstates+bx+0] ; ax of new task
064             mov bx, [cs:taskstates+bx+2] ; bx of new task
065             iret                 ; return to new task
066
067 start:      mov word [taskstates+10+4], taskone ; initialize ip
068             mov [taskstates+10+6], cs          ; initialize cs
069             mov word [taskstates+10+8], 0x0200 ; initialize flags
070             mov word [taskstates+20+4], tasktwo ; initialize ip
071             mov [taskstates+20+6], cs          ; initialize cs
072             mov word [taskstates+20+8], 0x0200 ; initialize flags
073             mov word [current], 0             ; set current task index
074
075             xor ax, ax
076             mov es, ax                       ; point es to IVT base
077             cli
078             mov word [es:8*4], timer
079             mov [es:8*4+2], cs                ; hook timer interrupt
080             mov ax, 0xb800
081             mov es, ax                       ; point es to video base
082             xor bx, bx                       ; initialize bx for tasks
083             sti
084
085             jmp $                             ; infinite loop

```

The space where all registers of a task are stored is called the process control block or PCB. Actual PCB contains a few more things that are not

relevant to us now. INT 08 that is saving and restoring the registers is called the scheduler and the whole event is called a context switch.

4.2. ELABORATE MULTITASKING

In our next example we will save all 14 registers and the stack as well. 28 bytes are needed by these registers in the PCB. We add some more space to make the size 32, a power of 2 for easy calculations. One of these words is used to form a linked list of the PCBs so that strict ordering of active PCBs is not necessary. Also in this example we have given every thread its own stack. Now threads can have function calls, parameters and local variables etc. Another important change in this example is that the creation of threads is now dynamic. The thread registration code initializes the PCB, and adds it to the linked list so that the scheduler will give it a turn.

Example 11.2

```

001 ; multitasking and dynamic thread registration
002 [org 0x0100]
003     jmp start
004
005 ; PCB layout:
006 ; ax,bx,cx,dx,si,di,bp,sp,ip,cs,ds,ss,es,flags,next,dummy
007 ; 0, 2, 4, 6, 8,10,12,14,16,18,20,22,24, 26, 28, 30
008
009 pcb:         times 32*16 dw 0           ; space for 32 PCBs
010 stack:       times 32*256 dw 0         ; space for 32 512 byte stacks
011 nextpcb:     dw 1                     ; index of next free pcb
012 current:     dw 0                     ; index of current pcb
013 lineno:      dw 0                     ; line number for next thread
014-057
058 ;;;; COPY LINES 028-071 FROM EXAMPLE 10.1 (printnum) ;;;;
059
060 ; mytask subroutine to be run as a thread
061 ; takes line number as parameter
062 mytask:      push bp
063             mov bp, sp
064             sub sp, 2                 ; thread local variable
065             push ax
066             push bx
067
068             mov ax, [bp+4]            ; load line number parameter
069             mov bx, 70                ; use column number 70
070             mov word [bp-2], 0        ; initialize local variable
071
072 printagain:  push ax                  ; line number
073             push bx                  ; column number
074             push word [bp-2]         ; number to be printed
075             call printnum            ; print the number
076             inc word [bp-2]          ; increment the local variable
077             jmp printagain           ; infinitely print
078
079             pop bx
080             pop ax
081             mov sp, bp
082             pop bp
083             ret
084
085 ; subroutine to register a new thread
086 ; takes the segment, offset, of the thread routine and a parameter
087 ; for the target thread subroutine
088 initpcb:    push bp
089             mov bp, sp
090             push ax
091             push bx
092             push cx
093             push si
094
095             mov bx, [nextpcb]         ; read next available pcb index
096             cmp bx, 32                ; are all PCBs used
097             je exit                   ; yes, exit
098

```

```

099      mov     cl, 5
100      shl     bx, cl           ; multiply by 32 for pcb start
101
102      mov     ax, [bp+8]       ; read segment parameter
103      mov     [pcb+bx+18], ax  ; save in pcb space for cs
104      mov     ax, [bp+6]       ; read offset parameter
105      mov     [pcb+bx+16], ax  ; save in pcb space for ip
106
107      mov     [pcb+bx+22], ds   ; set stack to our segment
108      mov     si, [nextpcb]    ; read this pcb index
109      mov     cl, 9
110      shl     si, cl           ; multiply by 512
111      add     si, 256*2+stack  ; end of stack for this thread
112      mov     ax, [bp+4]       ; read parameter for subroutine
113      sub     si, 2            ; decrement thread stack pointer
114      mov     [si], ax         ; pushing param on thread stack
115      sub     si, 2            ; space for return address
116      mov     [pcb+bx+14], si   ; save si in pcb space for sp
117
118      mov     word [pcb+bx+26], 0x0200 ; initialize thread flags
119      mov     ax, [pcb+28]      ; read next of 0th thread in ax
120      mov     [pcb+bx+28], ax   ; set as next of new thread
121      mov     ax, [nextpcb]    ; read new thread index
122      mov     [pcb+28], ax      ; set as next of 0th thread
123      inc     word [nextpcb]    ; this pcb is now used
124
125
126      exit:    pop     si
127              pop     cx
128              pop     bx
129              pop     ax
130              pop     bp
131              ret     6
132
133      ; timer interrupt service routine
134      timer:   push    ds
135              push    bx
136
137              push    cs
138              pop     ds       ; initialize ds to data segment
139
140              mov     bx, [current] ; read index of current in bx
141              shl     bx, 1
142              shl     bx, 1
143              shl     bx, 1
144              shl     bx, 1
145              shl     bx, 1     ; multiply by 32 for pcb start
146              mov     [pcb+bx+0], ax ; save ax in current pcb
147              mov     [pcb+bx+4], cx ; save cx in current pcb
148              mov     [pcb+bx+6], dx ; save dx in current pcb
149              mov     [pcb+bx+8], si ; save si in current pcb
150              mov     [pcb+bx+10], di ; save di in current pcb
151              mov     [pcb+bx+12], bp ; save bp in current pcb
152              mov     [pcb+bx+24], es ; save es in current pcb
153
154              pop     ax         ; read original bx from stack
155              mov     [pcb+bx+2], ax ; save bx in current pcb
156              pop     ax         ; read original ds from stack
157              mov     [pcb+bx+20], ax ; save ds in current pcb
158              pop     ax         ; read original ip from stack
159              mov     [pcb+bx+16], ax ; save ip in current pcb
160              pop     ax         ; read original cs from stack
161              mov     [pcb+bx+18], ax ; save cs in current pcb
162              pop     ax         ; read original flags from stack
163              mov     [pcb+bx+26], ax ; save cs in current pcb
164              mov     [pcb+bx+22], ss ; save ss in current pcb
165              mov     [pcb+bx+14], sp ; save sp in current pcb
166
167              mov     bx, [pcb+bx+28] ; read next pcb of this pcb
168              mov     [current], bx   ; update current to new pcb
169              mov     cl, 5
170              shl     bx, cl           ; multiply by 32 for pcb start
171
172              mov     cx, [pcb+bx+4] ; read cx of new process
173              mov     dx, [pcb+bx+6] ; read dx of new process
174              mov     si, [pcb+bx+8] ; read si of new process
175              mov     di, [pcb+bx+10] ; read diof new process

```

176		mov bp, [pcb+bx+12]	; read bp of new process
177		mov es, [pcb+bx+24]	; read es of new process
178		mov ss, [pcb+bx+22]	; read ss of new process
179		mov sp, [pcb+bx+14]	; read sp of new process
180			
181		push word [pcb+bx+26]	; push flags of new process
182		push word [pcb+bx+18]	; push cs of new process
183		push word [pcb+bx+16]	; push ip of new process
184		push word [pcb+bx+20]	; push ds of new process
185			
186		mov al, 0x20	
187		out 0x20, al	; send EOI to PIC
188			
189		mov ax, [pcb+bx+0]	; read ax of new process
190		mov bx, [pcb+bx+2]	; read bx of new process
191		pop ds	; read ds of new process
192		iret	; return to new process
193			
194	start:	xor ax, ax	
195		mov es, ax	; point es to IVT base
196			
197		cli	
198		mov word [es:8*4], timer	
199		mov [es:8*4+2], cs	; hook timer interrupt
200		sti	
201			
202	nextkey:	xor ah, ah	; service 0 - get keystroke
203		int 0x16	; bios keyboard services
204			
205		push cs	; use current code segment
206		mov ax, mytask	
207		push ax	; use mytask as offset
208		push word [lineno]	; thread parameter
209		call initpcb	; register the thread
210			
211		inc word [lineno]	; update line number
212		jmp nextkey	; wait for next keypress

When the program is executed the threads display the numbers independently. However as keys are pressed and new threads are registered, there is an obvious slowdown in the speed of multitasking. To improve that, we can change the timer interrupt frequency. The following can be used to set to an approximately 1ms interval.

```
mov ax, 1100
out 0x40, al
mov al, ah
out 0x40, al
```

This makes the threads look faster. However the only real change is that the timer interrupt is now coming more frequently.

4.3. MULTITASKING KERNEL AS TSR

The above examples had the multitasking code and the multithreaded code in one program. Now we separate the multitasking kernel into a TSR so that it becomes an operation system extension. We hook a software interrupt for the purpose of registering a new thread.

Example 11.3	
001	; multitasking kernel as a TSR
002	[org 0x0100]
003	jmp start
004	
005	; PCB layout:
006	; ax,bx,cx,dx,si,di,bp,sp,ip,cs,ds,ss,es,flags,next,dummy
007	; 0, 2, 4, 6, 8,10,12,14,16,18,20,22,24, 26, 28, 30
008	
009	pcb: times 32*16 dw 0 ; space for 32 PCBs

```

010      stack:      times 32*256 dw 0      ; space for 32 512 byte stacks
011      nextpcb:    dw 1                  ; index of next free pcb
012      current:    dw 0                  ; index of current pcb
013
014-073  ;;;; COPY LINES 133-192 FROM EXAMPLE 11.2 (timer) ;;;;
074
075      ; software interrupt to register a new thread
076      ; takes parameter block in ds:si
077      ; parameter block has cs, ip, ds, es, and param in this order
078      initpcb:     push ax
079                  push bx
080                  push cx
081                  push di
082
083                  mov bx, [cs:nextpcb]    ; read next available pcb index
084                  cmp bx, 32              ; are all PCBs used
085                  je exit                 ; yes, exit
086
087                  mov cl, 5
088                  shl bx, cl              ; multiply by 32 for pcb start
089
090                  mov ax, [si+0]          ; read code segment parameter
091                  mov [cs:pcb+bx+18], ax  ; save in pcb space for cs
092                  mov ax, [si+2]          ; read offset parameter
093                  mov [cs:pcb+bx+16], ax  ; save in pcb space for ip
094                  mov ax, [si+4]          ; read data segment parameter
095                  mov [cs:pcb+bx+20], ax  ; save in pcb space for ds
096                  mov ax, [si+6]          ; read extra segment parameter
097                  mov [cs:pcb+bx+24], ax  ; save in pcb space for es
098
099                  mov [cs:pcb+bx+22], cs  ; set stack to our segment
100                  mov di, [cs:nextpcb]    ; read this pcb index
101                  mov cl, 9
102                  shl di, cl              ; multiply by 512
103                  add di, 256*2+stack     ; end of stack for this thread
104                  mov ax, [si+8]          ; read parameter for subroutine
105                  sub di, 2               ; decrement thread stack pointer
106                  mov [cs:di], ax         ; pushing param on thread stack
107                  sub di, 4               ; space for far return address
108                  mov [cs:pcb+bx+14], di   ; save di in pcb space for sp
109
110                  mov word [cs:pcb+bx+26], 0x0200 ; initialize flags
111                  mov ax, [cs:pcb+28]     ; read next of 0th thread in ax
112                  mov [cs:pcb+bx+28], ax  ; set as next of new thread
113                  mov ax, [cs:nextpcb]    ; read new thread index
114                  mov [cs:pcb+28], ax     ; set as next of 0th thread
115                  inc word [cs:nextpcb]   ; this pcb is now used
116
117      exit:        pop di
118                  pop cx
119                  pop bx
120                  pop ax
121                  iret
122
123      start:       xor ax, ax
124                  mov es, ax              ; point es to IVT base
125
126                  mov word [es:0x80*4], initpcb
127                  mov [es:0x80*4+2], cs   ; hook software int 80
128                  cli
129                  mov word [es:0x08*4], timer
130                  mov [es:0x08*4+2], cs   ; hook timer interrupt
131                  sti
132
133                  mov dx, start
134                  add dx, 15
135                  mov cl, 4
136                  shr dx, cl
137
138                  mov ax, 0x3100          ; terminate and stay resident
139                  int 0x21
140

```

The second part of our example is a simple program that has the threads to be registered with the multitasking kernel using its exported services.

Example 11.4

```

001 ; multitasking TSR caller
002 [org 0x0100]
003     jmp start
004
005 ; parameter block layout:
006 ; cs,ip,ds,es,param
007 ; 0, 2, 4, 6, 8
008
009 paramblock: times 5 dw 0 ; space for parameters
010 lineno: dw 0 ; line number for next thread
011
012-055 ;;;; COPY LINES 028-071 FROM EXAMPLE 10.1 (printnum) ;;;;
056
057 ; subroutine to be run as a thread
058 ; takes line number as parameter
059 mytask: push bp
060         mov bp, sp
061         sub sp, 2 ; thread local variable
062         push ax
063         push bx
064
065         mov ax, [bp+4] ; load line number parameter
066         mov bx, 70 ; use column number 70
067         mov word [bp-2], 0 ; initialize local variable
068
069 printagain: push ax ; line number
070            push bx ; column number
071            push word [bp-2] ; number to be printed
072            call printnum ; print the number
073            inc word [bp-2] ; increment the local variable
074            jmp printagain ; infinitely print
075
076         pop bx
077         pop ax
078         mov sp, bp
079         pop bp
080         retf
081
082 start: mov ah, 0 ; service 0 - get keystroke
083        int 0x16 ; bios keyboard services
084
085        mov [paramblock+0], cs ; code segment parameter
086        mov word [paramblock+2], mytask ; offset parameter
087        mov [paramblock+4], ds ; data segment parameter
088        mov [paramblock+6], es ; extra segment parameter
089        mov ax, [lineno]
090        mov [paramblock+8], ax ; parameter for thread
091        mov si, paramblock ; address of param block in si
092        int 0x80 ; multitasking kernel interrupt
093
094        inc word [lineno] ; update line number
095        jmp start ; wait for next key

```

We introduce yet another use of the multitasking kernel with this new example. In this example three different sort of routines are multitasked by the same kernel instead of repeatedly registering the same routine.

Example 11.5

```

001 ; another multitasking TSR caller
002 [org 0x0100]
003     jmp start
004
005 ; parameter block layout:
006 ; cs,ip,ds,es,param
007 ; 0, 2, 4, 6, 8
008
009 paramblock: times 5 dw 0 ; space for parameters
010 lineno: dw 0 ; line number for next thread
011 chars: db '\|/-' ; chracters for rotating bar
012 message: db 'moving hello' ; moving string
013 message2: db ' ' ; to erase previous string

```

```

014      messagelen:   dw 12                      ; length of above strings
015
016-059  ;;;; COPY LINES 028-071 FROM EXAMPLE 10.1 (printnum) ;;;;
060-101  ;;;; COPY LINES 073-114 FROM EXAMPLE 10.1 (printstr) ;;;;
102
103      ; subroutine to run as first thread
104      mytask:       push bp
105                   mov  bp, sp
106                   sub  sp, 2                  ; thread local variable
107                   push ax
108                   push bx
109
110                   xor  ax, ax                  ; use line number 0
111                   mov  bx, 70                 ; use column number 70
112                   mov  word [bp-2], 0        ; initialize local variable
113
114      printagain:    push ax                   ; line number
115                   push bx                   ; column number
116                   push word [bp-2]          ; number to be printed
117                   call printnum             ; print the number
118                   inc  word [bp-2]          ; increment the local variable
119                   jmp  printagain           ; infinitely print
120
121                   pop  bx
122                   pop  ax
123                   mov  sp, bp
124                   pop  bp
125                   retf
126
127      ; subroutine to run as second thread
128      mytask2:      push ax
129                   push bx
130                   push es
131
132                   mov  ax, 0xb800
133                   mov  es, ax                ; point es to video base
134                   xor  bx, bx                ; initialize to use first shape
135
136      rotateagain:   mov  al, [chars+bx]      ; read current shape
137                   mov  [es:40], al          ; print at specified place
138                   inc  bx                   ; update to next shape
139                   and  bx, 3                 ; take modulus with 4
140                   jmp  rotateagain          ; repeat infinitely
141
142                   pop  es
143                   pop  bx
144                   pop  ax
145                   retf
146
147      ; subroutine to run as third thread
148      mytask3:      push bp
149                   mov  bp, sp
150                   sub  sp, 2                  ; thread local variable
151                   push ax
152                   push bx
153                   push cx
154
155                   mov  word [bp-2], 0        ; initialize line number to 0
156
157      nextline:      push word [bp-2]          ; line number
158                   mov  bx, 50
159                   push bx                   ; column number 50
160                   mov  ax, message
161                   push ax                   ; offset of string
162                   push word [messagelen]    ; length of string
163                   call printstr             ; print the string
164
165
166      waithere:      mov  cx, 0x100
167                   push cx                   ; save outer loop counter
168                   mov  cx, 0xffff
169                   loop $                     ; repeat ffff times
170                   pop  cx                   ; restore outer loop counter
171                   loop waithere             ; repeat 0x100 times
172
173                   push word [bp-2]          ; line number
174                   mov  bx, 50               ; column number 50

```



```

175      push bx
176      mov ax, message2
177      push ax
178      push word [messagelen] ; length of string
179      call printstr         ; print the string
180
181      inc word [bp-2]       ; update line number
182      cmp word [bp-2], 25   ; is this the last line
183      jne skipreset        ; no, proceed to draw
184      mov word [bp-2], 0    ; yes, reset line number to 0
185
186      skipreset: jmp nextline ; proceed with next drawing
187
188      pop cx
189      pop bx
190      pop ax
191      mov sp, bp
192      pop bp
193      retf
194
195      start:  mov [paramblock+0], cs ; code segment parameter
196              mov word [paramblock+2], mytask ; offset parameter
197              mov [paramblock+4], ds ; data segment parameter
198              mov [paramblock+6], es ; extra segment parameter
199              mov word [paramblock+8], 0 ; parameter for thread
200              mov si, paramblock      ; address of param block in si
201              int 0x80                ; multitasking kernel interrupt
202
203              mov [paramblock+0], cs ; code segment parameter
204              mov word [paramblock+2], mytask2 ; offset parameter
205              mov [paramblock+4], ds ; data segment parameter
206              mov [paramblock+6], es ; extra segment parameter
207              mov word [paramblock+8], 0 ; parameter for thread
208              mov si, paramblock      ; address of param block in si
209              int 0x80                ; multitasking kernel interrupt
210
211              mov [paramblock+0], cs ; code segment parameter
212              mov word [paramblock+2], mytask3 ; offset parameter
213              mov [paramblock+4], ds ; data segment parameter
214              mov [paramblock+6], es ; extra segment parameter
215              mov word [paramblock+8], 0 ; parameter for thread
216              mov si, paramblock      ; address of param block in si
217              int 0x80                ; multitasking kernel interrupt
218
219      jmp $

```

EXERCISES

1. Change the multitasking kernel such that a new two byte variable is introduced in the PCB. This variable contains the number of turns this process should be given. For example if the first PCB contains 20 in this variable, the switch to second process should occur after 20 timer interrupts (approx one second at default speed) and similarly the switch from second to third process should occur after the number given in the second process's PCB.
2. Change the scheduler of the multitasking kernel to enqueue the current process index a ready queue, and dequeue the next process index from it, and assign it to current. Therefore the next field of the PCB is no longer used. Use queue functions from Exercise 5.XX.
3. Add a function in the multitasking kernel to fork the current process through a software interrupt. Fork should allocate a new PCB and copy values of all registers of the caller's PCB to the new PCB. It should allocate a stack and change SS, SP appropriately in the new PCB. It has to copy the caller's stack on the newly allocated stack. It will set AX in the new PCB to 0 and in the old PB to 1 so that both threads can identify which is the creator and which is the created process and can act accordingly.

4. Add a function in the multitasking kernel accessible via a software interrupt that allows the current process to terminate itself.
5. Create a queue in the multitasking kernel called kbQ. This queue initially empty will contain characters typed by the user. Hook the keyboard interrupt for getting user keys. Convert the scan code to ASCII if the key is from a-z or 0-9 and enqueue it in kbQ. Ignore all other scan codes. Write a function checkkey accessible via a software interrupt that returns the process in AX a value removed from the queue. It waits if there is no key in the queue. Be aware of enabling interrupts if you wait here.
6. Modify the multitasking kernel such that the initial process displays at the last line of the screen whatever is typed by the user and clears that line on enter. If the user types quit followed by enter restore everything to normal as it was before the multitasking kernel was there. If the user types start followed by enter, start one more rotating bar on the screen. The first rotating bar should appear in the upper left, the next in the second column, then third and so on. The bar color should be white. The user can type the commands 'white', 'red', and 'green' to change the color of new bars.

5

Video Services

5.1. BIOS VIDEO SERVICES

The Basic Input Output System (BIOS) provides services for video, keyboard, serial port, parallel port, time etc. The video services are exported via INT 10. We will discuss some very simple services. Video services are classified into two broad categories; graphics mode services and text mode services. In graphics mode a location in video memory corresponds to a dot on the screen. In text mode this relation is not straightforward. The video memory holds the ASCII of the character to be shown and the actual shape is read from a font definition stored elsewhere in memory. We first present a list of common video services used in text mode.

INT 10 - VIDEO - SET VIDEO MODE

AH = 00h

AL = desired video mode

Some common video modes include 40x25 text mode (mode 0), 80x25 text mode (mode 2), 80x50 text mode (mode 3), and 320x200 graphics mode (mode D).

INT 10 - VIDEO - SET TEXT-MODE CURSOR SHAPE

AH = 01h

CH = cursor start and options

CL = bottom scan line containing cursor (bits 0-4)

INT 10 - VIDEO - SET CURSOR POSITION

AH = 02h

BH = page number

0-3 in modes 2&3

0-7 in modes 0&1

0 in graphics modes

DH = row (00h is top)

DL = column (00h is left)

INT 10 - VIDEO - SCROLL UP WINDOW

AH = 06h

AL = number of lines by which to scroll up (00h = clear entire window)

BH = attribute used to write blank lines at bottom of window

CH, CL = row, column of window's upper left corner

DH, DL = row, column of window's lower right corner

INT 10 - VIDEO - SCROLL DOWN WINDOW

AH = 07h

AL = number of lines by which to scroll down (00h=clear entire window)

BH = attribute used to write blank lines at top of window

CH, CL = row, column of window's upper left corner

DH, DL = row, column of window's lower right corner

INT 10 - VIDEO - WRITE CHARACTER AND ATTRIBUTE AT CURSOR POSITION

AH = 09h

AL = character to display

BH = page number

```

BL = attribute (text mode) or color (graphics mode)
CX = number of times to write character
INT 10 - VIDEO - WRITE CHARACTER ONLY AT CURSOR POSITION
AH = 0Ah
AL = character to display
BH = page number
BL = attribute (text mode) or color (graphics mode)
CX = number of times to write character
INT 10 - VIDEO - WRITE STRING
AH = 13h
AL = write mode
    bit 0: update cursor after writing
    bit 1: string contains alternating characters and attributes
    bits 2-7: reserved (0)
BH = page number
BL = attribute if string contains only characters
CX = number of characters in string
DH, DL = row, column at which to start writing
ES:BP -> string to write

```


Chargen Services

In our first example we will read the font definition in memory and change it to include a set of all on pixels in the last line showing an effect of underline on all character including space. An 8x16 font is stored in 16 bytes. A sample character and the corresponding 16 values stored in the font information are shown for the character 'A'. We start with two services from the chargen subset of video services that we are going to use.

```

INT 10 - VIDEO - GET FONT INFORMATION
AX = 1130h
BH = pointer specifier
Return:
ES:BP = specified pointer
CX = bytes/character of on-screen font
DL = highest character row on screen
INT 10 - TEXT-MODE CHARGEN
AX = 1110h
ES:BP -> user table
CX = count of patterns to store
DX = character offset into map 2 block
BL = block to load in map 2
BH = number of bytes per character pattern

```



We will use 6 as the pointer specifier which means the 8x16 font stored in ROM.

Example 12.1			
001	; put underlines on screen font		
002	[org 0x0100]		
003		jmp	start
004			
005	font:	times 256*16 db 0	; space for font
006			
007	start:	mov ax, 0x1130	; service 11/30 - get font info
008		mov bx, 0x0600	; ROM 8x16 font
009		int 0x10	; bios video services
010			
011		mov si, bp	; point si to rom font data
012		mov di, font	; point di to space for font

```

013      mov cx, 256*16      ; font size
014      push ds
015      push es
016      pop ds              ; ds:si to rom font data
017      pop es              ; es:di to space for font
018      cld                 ; auto increment mode
019      rep movsb           ; copy font
020
021      push cs
022      pop ds              ; restore ds to data segment
023
024      mov si, font-1      ; point si before first char
025      mov cx, 0x100       ; total 256 characters
026      change: add si, 16   ; one character has 16 bytes
027      mov byte [si], 0xFF ; change last line to all ones
028      loop change        ; repeat for each character
029
030      mov bp, font        ; es:bp points to new font
031      mov bx, 0x1000      ; bytes per char & block number
032      mov cx, 0x100       ; number of characters to change
033      xor dx, dx          ; first character to change
034      mov ax, 0x1110      ; service 11/10 - load user font
035      int 0x10            ; bios video services
036
037      mov ax, 0x4c00       ; terminate program
038      int 0x21

```

Our second example is similar to the last example however in this case we are doing something funny on the screen. We are reversing the shapes of all the characters on the screen.

Example 12.2

```

001      ; reverse each character of screen font
002      [org 0x0100]
003      jmp start
004
005      font: times 256*16 db 0      ; space for font
006
007      start: mov ax, 0x1130        ; service 11/30 - get font info
008      mov bx, 0x0600              ; ROM 8x16 font
009      int 0x10                    ; bios video services
010
011      mov si, bp                  ; point si to rom font data
012      mov di, font                ; point di to space for font
013      mov cx, 256*16              ; font size
014      push ds
015      push es
016      pop ds                      ; ds:si to rom font data
017      pop es                      ; es:di to space for font
018      cld                         ; auto increment mode
019      rep movsb                   ; copy font
020
021      push cs
022      pop ds                      ; restore ds to data segment
023
024      mov si, font                ; point si to start of font
025      change: mov al, [si]         ; read one byte
026      mov cx, 8
027      inner: shl al, 1             ; shift left with MSB in carry
028      rcr bl, 1                   ; rotate right using carry
029      loop inner                  ; repeat eight times
030      mov [si], bl                ; write back reversed byte
031      inc si                      ; next byte of font
032      cmp si, font+256*16         ; is whole font reversed
033      jne change                  ; no, reverse next byte
034
035      mov bp, font                ; es:bp points to new font
036      mov bx, 0x1000              ; bytes per char & block number
037      mov cx, 0x100               ; number of characters to change
038      xor dx, dx                  ; first character to change
039      mov ax, 0x1110              ; service 11/10 - load user font
040      int 0x10                    ; bios video services
041

```

042	mov ax, 0x4c00	; terminate program
043	int 0x21	

Graphics Mode Services

We will take an example of using graphics mode video services as well. We will draw a line across the screen using the following service.

```

INT 10 - VIDEO - WRITE GRAPHICS PIXEL
AH = 0Ch
BH = page number
AL = pixel color
CX = column
DX = row

```

Example 12.3

001	; draw line in graphics mode	
002	[org 0x0100]	
003	mov ax, 0x000D	; set 320x200 graphics mode
004	int 0x10	; bios video services
005		
006	mov ax, 0x0C07	; put pixel in white color
007	xor bx, bx	; page number 0
008	mov cx, 200	; x position 200
009	mov dx, 200	; y position 200
010		
011	11: int 0x10	; bios video services
012	dec dx	; decrease y position
013	loop 11	; decrease x position and repeat
014		
015	mov ah, 0	; service 0 - get keystroke
016	int 0x16	; bios keyboard services
017		
018	mov ax, 0x0003	; 80x25 text mode
019	int 0x10	; bios video services
020		
021	mov ax, 0x4c00	; terminate program
022	int 0x21	

5.2. DOS VIDEO SERVICES

Services of DOS are more cooked and at a higher level than BIOS. They provide less control but make routine tasks much easier. Some important DOS services are listed below.

```

INT 21 - READ CHARACTER FROM STANDARD INPUT, WITH ECHO
AH = 01h
Return: AL = character read
INT 21 - WRITE STRING TO STANDARD OUTPUT
AH = 09h
DS:DX -> $ terminated string
INT 21 - BUFFERED INPUT
AH = 0Ah
DS:DX -> dos input buffer

```

The DOS input buffer has a special format where the first byte stores the maximum characters buffer can hold, the second byte holds the number of characters actually read on return, and the following space is used for the actual characters read. We start with an example of reading a string with service 1 and displaying it with service 9.

Example 12.4

```

001 ; character input using dos services
002 [org 0x0100]
003     jmp start
004
005 maxlength: dw 80 ; maximum length of input
006 message: db 10, 13, 'hello $' ; greetings message
007 buffer: times 81 db 0 ; space for input string
008
009 start: mov cx, [maxlength] ; load maximum length in cx
010     mov si, buffer ; point si to start of buffer
011
012 nextchar: mov ah, 1 ; service 1 - read character
013     int 0x21 ; dos services
014
015     cmp al, 13 ; is enter pressed
016     je exit ; yes, leave input
017     mov [si], al ; no, save this character
018     inc si ; increment buffer pointer
019     loop nextchar ; repeat for next input char
020
021 exit: mov byte [si], '$' ; append $ to user input
022
023     mov dx, message ; greetings message
024     mov ah, 9 ; service 9 - write string
025     int 0x21 ; dos services
026
027     mov dx, buffer ; user input buffer
028     mov ah, 9 ; service 9 - write string
029     int 0x21 ; dos services
030
031     mov ax, 0x4c00 ; terminate program
032     int 0x21

```

Our next example uses the more cooked buffered input service of DOS and using the same service 9 to print the string.

Example 12.5

```

001 ; buffer input using dos services
002 [org 0x0100]
003     jmp start
004
005 message: db 10,13,'hello ', 10, 13, '$'
006 buffer: db 80 ; length of buffer
007     db 0 ; number of character on return
008     times 80 db 0 ; actual buffer space
009
010 start: mov dx, buffer ; input buffer
011     mov ah, 0x0A ; service A - buffered input
012     int 0x21 ; dos services
013
014     mov bh, 0
015     mov bl, [buffer+1] ; read actual size in bx
016     mov byte [buffer+2+bx], '$' ; append $ to user input
017
018     mov dx, message ; greetings message
019     mov ah, 9 ; service 9 - write string
020     int 0x21 ; dos services
021
022     mov dx, buffer+2 ; user input buffer
023     mov ah, 9 ; service 9 - write string
024     int 0x21 ; dos services
025
026     mov ax, 0x4c00 ; terminate program
027     int 0x21

```

More detail of DOS and BIOS interrupts is available in the Ralph Brown Interrupt List.


6

Secondary Storage

6.1. PHYSICAL FORMATION

A floppy disk is a circular plate with a fine coating of magnetic material over it. The plate is enclosed in a plastic jacket which has a cover that can slide to expose the magnetic surface. The drive motor attaches itself to the central piece and rotates the plate. Two heads on both sides can read the magnetically encoded data on the disk.

If the head is fixed and the motor rotates the disk the readable area on the disk surface forms a circle called a track. Head moved to the next step forms another track and so on. In hard disks the same structure is extended to a larger number of tracks and plates. The tracks are further cut vertically into sectors. This is a logical division of the area on the tracks. Each sector holds 512 bytes of data. A standard floppy disk has 80 tracks and 18 sectors per track with two heads, one on each side totalling to 2880 sectors or 1440 KB of data. Hard disks have varying number of heads and tracks pertaining to their different capacities.



BIOS sees the disks as a combination of sectors, tracks, and heads, as a raw storage device without concern to whether it is reading a file or directory. BIOS provides the simplest and most powerful interface to the storage medium. However this raw storage is meaningless to the user who needs to store his files and organize them into directories. DOS builds a logical structure on this raw storage space to provide these abstractions. This logical formation is read and interpreted by DOS. If another file system is build on the same storage medium the interpretations change. Main units of the DOS structure are the boot sector in head 0, track 0, and sector 1, the first FAT starting from head 0, track 0, sector 2, the second copy of FAT starting from head 0, track 0, sector 11, and the root directory starting from head 1, track 1, sector 2. The area from head 0, track 1, sector 16 to head 1, track 79, sector 18 is used for storing the data of the files. Among this we will be exploring the directory structure further. The 32 sectors reserved for the root directory contain 512 directory entries. The format of a 32 byte directory entry is shown below.

```
+00 Filename (8 bytes)
+08 Extension (3 bytes)
+0B Flag Byte (1 byte)
+0C Reserved (1 byte)
+0D Creation Date/Time (5 bytes)
+12 Last Accessed Data (2 bytes)
+14 Starting Cluster High Word (2 bytes) for FAT32
+16 Time (2 bytes)
+18 Date (2 bytes)
+1A Starting Cluster Low Word (2 bytes)
+1C File Size (4 bytes)
```

6.2. STORAGE ACCESS USING BIOS

We will be using BIOS disk services to directly see the data stored in the directory entries by DOS. For this purpose we will be using the BIOS disk services.

```
INT 13 - DISK - RESET DISK SYSTEM
AH = 00h
DL = drive
Return:
CF = error flag
AH = error code

INT 13 - DISK - READ SECTOR(S) INTO MEMORY
AH = 02h
AL = number of sectors to read (must be nonzero)
CH = low eight bits of cylinder number
CL = sector number 1-63 (bits 0-5)
    high two bits of cylinder (bits 6-7, hard disk only)
DH = head number
DL = drive number (bit 7 set for hard disk)
ES:BX -> data buffer
Return:
CF = error flag
AH = error code
AL = number of sectors transferred

INT 13 - DISK - WRITE DISK SECTOR(S)
AH = 03h
AL = number of sectors to write (must be nonzero)
CH = low eight bits of cylinder number
CL = sector number 1-63 (bits 0-5)
    high two bits of cylinder (bits 6-7, hard disk only)
DH = head number
DL = drive number (bit 7 set for hard disk)
ES:BX -> data buffer
Return:
CF = error flag
AH = error code
AL = number of sectors transferred

INT 13 - DISK - GET DRIVE PARAMETERS
AH = 08h
DL = drive (bit 7 set for hard disk)
Return:
CF = error flag
AH = error code
CH = low eight bits of maximum cylinder number
CL = maximum sector number (bits 5-0)
```

high two bits of maximum cylinder number (bits 7-6)
 DH = maximum head number
 DL = number of drives
 ES:DI -> drive parameter table (floppies only)

Example 13.1

```

001 ; floppy directory using bios services
002 [org 0x0100]
003     jmp start
004
005 sector: times 512 db 0 ; space for directory sector
006 entryname: times 11 db 0 ; space for a file name
007     db 10, 13, '$' ; new line and terminating $
008
009 start: mov ah, 0 ; service 0 - reset disk system
010     mov dl, 0 ; drive A:
011     int 0x13 ; bios disk services
012     jc error ; if error, terminate program
013
014     mov ah, 2 ; service 2 - read sectors
015     mov al, 1 ; count of sectors
016     mov ch, 0 ; cylinder
017     mov cl, 2 ; sector
018     mov dh, 1 ; head
019     mov dl, 0 ; drive A:
020     mov bx, sector ; buffer to read sector
021     int 0x13 ; bios disk services
022     jc error ; if error, terminate program
023
024     mov bx, 0 ; start from first entry
025 nextentry: mov di, entryname ; point di to space for filename
026     mov si, sector ; point si to sector
027     add si, bx ; move ahead to desired entry
028     mov cx, 11 ; one filename is 11 bytes long
029     cld ; auto increment mode
030     rep movsb ; copy filename
031
032     mov ah, 9 ; service 9 - output string
033     mov dx, entryname ; filename to be printed
034     int 0x21 ; dos services
035
036     add bx, 32 ; point to next dir entry
037     cmp bx, 512 ; is last entry in this sector
038     jne nextentry ; no, print next entry
039
040 error: mov ax, 0x4c00 ; terminate program
041     int 0x21

```

With the given services and the bits allocated for heads, tracks, and sectors only 8GB disks can be accessed. This limitation can be overcome by using INT 13 extensions that take a linear 64bit sector number and handle all the head, track, sector conversion themselves. The important services in this category are listed below.

INT 13 - INT 13 Extensions - EXTENDED READ

AH = 42h
 DL = drive number
 DS:SI -> disk address packet
 Return:
 CF = error flag
 AH = error code
 disk address packet's block count field set to number of blocks successfully transferred

INT 13 - INT 13 Extensions - EXTENDED WRITE

AH = 43h
 AL = write flags

```

DL = drive number
DS:SI -> disk address packet
Return:
CF = error flag
AH = error code
        disk address packet's block count field set to number of blocks
        successfully transferred

```

The format of the disk address packet used above is as follows.

Offset	Size	Description
00h	BYTE	size of packet = 10h
01h	BYTE	reserved (0)
02h	WORD	number of blocks to transfer
04h	DWORD	-> transfer buffer
08h	QWORD	starting absolute block number

Hard disks have a different formation from floppy disks in that there is a partition table at the start that allows several logical disks to be maintained within a single physical disk. The physical sector 0 holds the master boot record and a partition table towards the end. The first 446 bytes contain MBR, then there are 4 16 byte partition entries and then there is a 2 byte signature. A partition table entry has the following format.

```

Byte 0 - 0x80 for active 0x00 for inactive
Byte 1-3 - Starting CHS
Byte 4 - Partition Type
Byte 5-7 - Ending CHS
Byte 8-B - Starting LBA
Byte C-F - Size of Partition

```

Some important partition types are listed below.

```

00 Unused Entry
01 FAT12
05 Extended Partition
06 FAT16
0b FAT32
0c FAT32 LBA
0e FAT16 LBA
0f Extended LBA
07 NTFS

```

Extended partition type signals that the specified area is treated as a complete hard disk with its own partition table and partitions. Therefore extended partitions allow a recursion in partitioning and consequently an infinite number of partitions are possible. The following program reads the partition tables (primary and extended) using recursion and displays in an indented form all partitions present on the first hard disk in the system.

Example 13.2	
001	; a program to display the partition table
002	[org 0x0100]
003	jmp start
004	
005	dap: db 0x10, 0 ; disk address packet
006	dw 1
007	dd 0, 0, 0

```

008
009-026 msg:          times 17 db ' '
027                  db 10, 13, '$'
028 fat12:            db 'FAT12...$'
029 fat16:            db 'FAT16...$'
030 fat32:            db 'FAT32...$'
031 ntfs:             db 'NTFS...$'
032 extended:        db 'EXTEND...$'
033 unknown:         db 'UNKNOWN.$'
034
035 partypes:         dw 0x1, fat12      ; table of known partition types
036                  dw 0x5, extended
037                  dw 0x6, fat16
038                  dw 0xe, fat16
039                  dw 0xb, fat32
040                  dw 0xc, fat32
041                  dw 0x7, ntfs
042                  dw 0xf, extended
043                  dw 0x0, unknown
044
045 ; subroutine to print a number in a string as hex
046 ; takes address of string and a 16bit number as parameter
047 printnum:         push bp
048                  mov bp, sp
049                  push ax
050                  push bx
051                  push cx
052                  push dx
053                  push di
054
055                  mov di, [bp+6]      ; string to store the number
056                  add di, 3
057
058                  mov ax, [bp+4]      ; load number in ax
059                  mov bx, 16          ; use base 16 for division
060                  mov cx, 4
061
062 nextdigit:        mov dx, 0
063                  div bx              ; divide by 16
064                  add dl, 0x30        ; convert into ascii value
065                  cmp dl, 0x39
066                  jbe skipalpha
067
068                  add dl, 7
069
070 skipalpha:        mov [di], dl       ; update char in string
071                  dec di
072                  loop nextdigit
073
074                  pop di
075                  pop dx
076                  pop cx
077                  pop bx
078                  pop ax
079                  pop bp
080                  ret 4
081
082 ; subroutine to print the start and end of a partition
083 ; takes the segment and offset of the partition table entry
084 printpart:        push bp
085                  mov bp, sp
086                  push es
087                  push ax
088                  push di
089
090                  les di, [bp+4]      ; point es:di to dap
091
092                  mov ax, msg
093                  push ax
094                  push word [es:di+0xA]
095                  call printnum      ; print first half of start
096
097                  add ax, 4
098                  push ax
099                  push word [es:di+0x8]
100                  call printnum      ; print second half of start

```

```

101
102         add ax, 5
103         push ax
104         push word [es:di+0xE]
105         call printnum           ; print first half of end
106
107         add ax, 4
108         push ax
109         push word [es:di+0xC]
110         call printnum           ; print second half of end
111
112         mov dx, msg
113         mov ah, 9
114         int 0x21                ; print the whole on the screen
115
116         pop di
117         pop ax
118         pop es
119         pop bp
120         ret 4
121
122         ; recursive subroutine to read the partition table
123         ; take indentation level and 32bit absolute block number as parameters
124 readpart: push bp
125         mov bp, sp
126         sub sp, 512             ; local space to read sector
127         push ax
128         push bx
129         push cx
130         push dx
131         push si
132
133         mov ax, bp
134         sub ax, 512
135         mov word [dap+4], ax    ; init dest offset in dap
136         mov [dap+6], ds        ; init dest segment in dap
137         mov ax, [bp+4]
138         mov [dap+0x8], ax      ; init sector no in dap
139         mov ax, [bp+6]
140         mov [dap+0xA], ax      ; init second half of sector no
141
142         mov ah, 0x42            ; read sector in LBA mode
143         mov dl, 0x80            ; first hard disk
144         mov si, dap             ; address of dap
145         int 0x13                ; int 13
146
147         jc failed               ; if failed, leave
148
149         mov si, -66             ; start of partition info
150 nextpart: mov ax, [bp+4]         ; read relative sector number
151         add [bp+si+0x8], ax      ; make it absolute
152         mov ax, [bp+6]         ; read second half
153         adc [bp+si+0xA], ax      ; make seconf half absolute
154
155         cmp byte [bp+si+4], 0    ; is partition unused
156         je exit
157
158         mov bx, partypes        ; point to partition types
159         mov di, 0
160 nextmatch: mov ax, [bx+di]
161         cmp [bp+si+4], al        ; is this partition known
162         je found                 ; yes, so print its name
163         add di, 4                ; no, try next entry in table
164         cmp di, 32               ; are all entries compared
165         jne nextmatch            ; no, try another
166
167         found: mov cx, [bp+8]    ; load indentation level
168         jcxz noindent            ; skip if no indentation needed
169 indent:  mov dl, ' '
170         mov ah, 2                ; display char service
171         int 0x21                ; dos services
172         loop indent              ; print required no of spaces
173
174         noindent: add di, 2
175         mov dx, [bx+di]          ; point to partition type name
176         mov ah, 9                ; print string service
177

```

```

178          int 0x21          ; dos services
179
180          push ss
181          mov ax, bp
182          add ax, si
183          push ax            ; pass partition entry address
184          call printpart     ; print start and end from it
185
186          cmp byte [bp+si+4], 5 ; is it an extended partition
187          je recurse        ; yes, make a recursive call
188
189          cmp byte [bp+si+4], 0xf ; is it an extended partition
190          jne exit          ; yes, make a recursive call
191
192  recurse:  mov ax, [bp+8]
193            add ax, 2        ; increase indentation level
194            push ax
195            push word [bp+si+0xA] ; push partition type address
196            push word [bp+si+0x8]
197            call readpart    ; recursive call
198
199  exit:     add si, 16        ; point to next partition entry
200            cmp si, -2       ; gone past last entry
201            jne nextpart     ; no, read this entry
202
203  failed:   pop si
204            pop dx
205            pop bx
206            pop cx
207            pop ax
208            mov sp, bp
209            pop bp
210            ret 6
211
212  start:    xor ax, ax
213            push ax          ; start from zero indentation
214            push ax          ; main partition table at 0
215            push ax
216            call readpart    ; read and print it
217
218            mov ax, 0x4c00    ; terminate program
219            int 0x21

```

6.3. STORAGE ACCESS USING DOS

BIOS provides raw access to the storage medium while DOS gives a more logical view and more cooked services. Everything is a file. A directory is a specially organized file that is interpreted by the operating system itself. A list of important DOS services for file manipulation is given below.

INT 21 - CREATE OR TRUNCATE FILE

AH = 3Ch
CX = file attributes
DS:DX -> ASCIZ filename
Return:
CF = error flag
AX = file handle or error code

INT 21 - OPEN EXISTING FILE

AH = 3Dh
AL = access and sharing modes
DS:DX -> ASCIZ filename
CL = attribute mask of files to look for (server call only)
Return:
CF = error flag
AX = file handle or error code

INT 21 - CLOSE FILE

AH = 3Eh

```
BX = file handle
Return:
CF = error flag
AX = error code
INT 21 - READ FROM FILE
AH = 3Fh
BX = file handle
CX = number of bytes to read
DS:DX -> buffer for data
Return:
CF = error flag
AX = number of bytes actually read or error code
INT 21 - WRITE TO FILE
AH = 40h
BX = file handle
CX = number of bytes to write
DS:DX -> data to write
Return:
CF = error flag
AX = number of bytes actually written or error code
INT 21 - DELETE FILE
AH = 41h
DS:DX -> ASCIZ filename (no wildcards, but see notes)
Return:
CF = error flag
AX = error code
INT 21 - SET CURRENT FILE POSITION
AH = 42h
AL = origin of move
BX = file handle
CX:DX = offset from origin of new file position
Return:
CF = error flag
DX:AX = new file position in bytes from start of file
AX = error code in case of error
INT 21 - GET FILE ATTRIBUTES
AX = 4300h
DS:DX -> ASCIZ filename
Return:
CF = error flag
CX = file attributes
AX = error code
INT 21 - SET FILE ATTRIBUTES
AX = 4301h
CX = new file attributes
DS:DX -> ASCIZ filename
Return:
CF = error flag
AX = error code
```

We will use some of these services to find that two files are same in contents or different. We will read the file names from the command prompt. The command string is passed to the program in the program segment prefix located at offset 0 in the current segment. The area from 0-7F contains information for DOS, while the command tail length is stored at 80. From 81 to FF, the actual command tail is stored terminated by a CR (Carriage Return).

Example 13.3

```

001 ; file comparison using dos services
002 [org 0x0100]
003         jmp  start
004
005 filename1:  times 128 db 0          ; space for first filename
006 filename2:  times 128 db 0          ; space for second filename
007 handle1:    dw  0                   ; handle for first file
008 handle2:    dw  0                   ; handle for second file
009 buffer1:    times 4096 db 0         ; buffer for first file
010 buffer2:    times 4096 db 0         ; buffer for second file
011
012 format:     db  'Usage error: diff <filename1> <filename2>$('
013 openfailed: db  'First file could not be opened$('
014 openfailed2: db 'Second file could not be opened$('
015 readfailed: db  'First file could not be read$('
016 readfailed2: db 'Second file could not be read$('
017 different:  db  'Files are different$('
018 same:       db  'Files are same$('
019
020 start:      mov  ch, 0
021            mov  cl, [0x80]           ; command tail length in cx
022            dec  cx                   ; leave the first space
023            mov  di, 0x82             ; start of command tail in di
024            mov  al, 0x20             ; space for parameter separation
025            cld                       ; auto increment mode
026            repne scasb               ; search space
027            je   param2               ; if found, proceed
028            mov  dx, format            ; else, select error message
029            jmp  error                ; proceed to error printing
030
031 param2:     push cx                   ; save original cx
032            mov  si, 0x82             ; set si to start of param
033            mov  cx, di               ; set di to end of param
034            sub  cx, 0x82             ; find param size in cx
035            dec  cx                   ; excluding the space
036            mov  di, filename1        ; set di to space for filename 1
037            rep  movsb                ; copy filename there
038            mov  byte [di], 0         ; terminate filename with 0
039            pop  cx                   ; restore original cx
040            inc  si                   ; go to start of next filename
041            mov  di, filename2        ; set di to space for filename 2
042            rep  movsb                ; copy filename there
043            mov  byte [di], 0         ; terminate filename with 0
044
045            mov  ah, 0x3d             ; service 3d - open file
046            mov  al, 0                ; readonly mode
047            mov  dx, filename1        ; address of filename
048            int  0x21                 ; dos services
049            jnc  open2                ; if no error, proceed
050            mov  dx, openfailed        ; else, select error message
051            jmp  error                ; proceed to error printing
052
053 open2:      mov  [handle1], ax        ; save handle for first file
054            mov  ah, 0x3d             ; service 3d - open file
055            mov  al, 0                ; readonly mode
056            mov  dx, filename2        ; address of filename
057            int  0x21                 ; dos services
058            jnc  store2                ; if no error, proceed
059            mov  dx, openfailed2       ; else, select error message
060            jmp  error                ; proceed to error printing
061
062 store2:     mov  [handle2], ax        ; save handle for second file
063
064 readloop:   mov  ah, 0x3f             ; service 3f - read file
065            mov  bx, [handle1]        ; handle for file to read
066            mov  cx, 4096             ; number of bytes to read
067            mov  dx, buffer1          ; buffer to read in
068            int  0x21                 ; dos services
069            jnc  read2                ; if no error, proceed
070            mov  dx, readfailed        ; else, select error message
071            jmp  error                ; proceed to error printing
072
073 read2:      push ax                   ; save number of bytes read
074            mov  ah, 0x3f             ; service 3f - read file

```

075		mov bx, [handle2]	; handle for file to read
076		mov cx, 4096	; number of bytes to read
077		mov dx, buffer2	; buffer to read in
078		int 0x21	; dos services
079		jnc check	; if no error, proceed
080		mov dx, readfailed2	; else, select error message
081		jmp error	; proceed to error printing
082			
083	check:	pop cx	; number of bytes read of file 1
084		cmp ax, cx	; are number of byte same
085		je check2	; yes, proceed to compare them
086		mov dx, different	; no, files are different
087		jmp error	; proceed to message printing
088			
089	check2:	test ax, ax	; are zero bytes read
090		jnz compare	; no, compare them
091		mov dx, same	; yes, files are same
092		jmp error	; proceed to message printing
093			
094	compare:	mov si, buffer1	; point si to file 1 buffer
095		mov di, buffer2	; point di to file 2 buffer
096		repe cmpsb	; compare the two buffers
097		je check3	; if equal, proceed
098		mov dx, different	; else, files are different
099		jmp error	; proceed to message printing
100			
101	check3:	cmp ax, 4096	; were 4096 bytes read
102		je readloop	; yes, try to read more
103		mov dx, same	; no, files are same
104			
105	error:	mov ah, 9	; service 9 - output message
106		int 0x21	; dos services
107			
108		mov ah, 0x3e	; service 3e - close file
109		mov bx, [handle1]	; handle of file to close
110		int 0x21	; dos services
111			
112		mov ah, 0x3e	; service 3e - close file
113		mov bx, [handle2]	; handle of file to close
114		int 0x21	; dos services
115			
116		mov ax, 0x4c00	; terminate program
117		int 0x21	

Another interesting service that DOS provides regarding files is executing them. An important point to understand here is that whenever a program is executed in DOS all available memory is allocated to it. No memory is available to execute any new programs. Therefore memory must be freed using explicit calls to DOS for this purpose before a program is executed. Important services in this regard are listed below.

INT 21 - ALLOCATE MEMORY

AH = 48h

BX = number of paragraphs to allocate

Return:

CF = error flag

AX = segment of allocated block or error code in case of error

BX = size of largest available block in case of error

INT 21 - FREE MEMORY

AH = 49h

ES = segment of block to free

Return:

CF = error flag

AX = error code

INT 21 - RESIZE MEMORY BLOCK

AH = 4Ah

BX = new size in paragraphs

ES = segment of block to resize

```

Return:
CF = error flag
AX = error code
BX = maximum paragraphs available for specified memory block
INT 21 - LOAD AND/OR EXECUTE PROGRAM
AH = 4Bh
AL = type of load (0 = load and execute)
DS:DX -> ASCIZ program name (must include extension)
ES:BX -> parameter block
Return:
CF = error flag
AX = error code

```

The format of parameter block is as follows.

Offset	Size	Description
00h	WORD	segment of environment to copy for child process (copy caller's environment if 0000h)
02h	DWORD	pointer to command tail to be copied into child's PSP
06h	DWORD	pointer to first FCB to be copied into child's PSP
0Ah	DWORD	pointer to second FCB to be copied into child's PSP
0Eh	DWORD	(AL=01h) will hold subprogram's initial SS:SP on return
12h	DWORD	(AL=01h) will hold entry point (CS:IP) on return

As an example we will use the multitasking kernel client from the multitasking chapter and modify it such that after running all three threads it executes a new instance of the command prompt instead of indefinitely hanging around.

Example 13.4

```

001 ; another multitasking TSR caller
002 [org 0x0100]
003 jmp start
004
005 ; parameter block layout:
006 ; cs,ip,ds,es,param
007 ; 0, 2, 4, 6, 8
008
009 paramblock: times 5 dw 0 ; space for parameters
010 lineno: dw 0 ; line number for next thread
011 chars: db '\|/-' ; chracters for rotating bar
012 message: db 'moving hello' ; moving string
013 message2: db ' ' ; to erase previous string
014 messagelen: dw 12 ; length of above strings
015 tail: db ' ',13
016 command: db 'COMMAND.COM', 0
017 execblock: times 11 dw 0
018
019-062 ;;;; COPY LINES 028-071 FROM EXAMPLE 10.1 (prntnum) ;;;;
063-104 ;;;; COPY LINES 073-114 FROM EXAMPLE 10.1 (prntstr) ;;;;
104-127 ;;;; COPY LINES 103-126 FROM EXAMPLE 11.5 (mytask) ;;;;
128-146 ;;;; COPY LINES 128-146 FROM EXAMPLE 11.5 (mytask2) ;;;;
147-192 ;;;; COPY LINES 148-193 FROM EXAMPLE 11.5 (mytask3) ;;;;
193
194 start: mov [paramblock+0], cs ; code segment parameter
195 mov word [paramblock+2], mytask ; offset parameter
196 mov [paramblock+4], ds ; data segment parameter
197 mov [paramblock+6], es ; extra segment parameter
198 mov word [paramblock+8], 0 ; parameter for thread
199 mov si, paramblock ; address of param block in si
200 int 0x80 ; multitasking kernel interrupt
201
202 mov [paramblock+0], cs ; code segment parameter
203 mov word [paramblock+2], mytask2 ; offset parameter
204 mov [paramblock+4], ds ; data segment parameter

```

```

205      mov [paramblock+6], es ; extra segment parameter
206      mov word [paramblock+8], 0 ; parameter for thread
207      mov si, paramblock      ; address of param block in si
208      int 0x80                ; multitasking kernel interrupt
209
210      mov [paramblock+0], cs ; code segment parameter
211      mov word [paramblock+2], mytask3 ; offset parameter
212      mov [paramblock+4], ds ; data segment parameter
213      mov [paramblock+6], es ; extra segment parameter
214      mov word [paramblock+8], 0 ; parameter for thread
215      mov si, paramblock      ; address of param block in si
216      int 0x80                ; multitasking kernel interrupt
217
218      mov ah, 0x4a             ; service 4a - resize memory
219      mov bx, end              ; end of memory retained
220      add bx, 15               ; rounding up
221      mov cl, 4
222      shr bx, cl               ; converting into paras
223      int 0x21                ; dos services
224
225      mov ah, 0x4b             ; service 4b - exec
226      mov al, 0                ; load and execute
227      mov dx, command          ; command to be executed
228      mov bx, execblock        ; address of execblock
229      mov word [bx+2], tail    ; offset of command tail
230      mov [bx+4], ds           ; segment of command tail
231      int 0x21                ; dos services
232
233      jmp $                    ; loop infinitely if returned
234  end:

```

6.4. DEVICE DRIVERS

Device drivers are operating system extensions that become part of the operating system and extend its services to new devices. Device drivers in DOS are very simple. They just have their services exposed through the file system interface.

Device driver file starts with a header containing a link to the next driver in the first four bytes followed by a device attribute word. The most important bit in the device attribute word is bit 15 which dictates if it is a character device or a block device. If the bit is zero the device is a character device and otherwise a block device. Next word in the header is the offset of a strategy routine, and then is the offset of the interrupt routine and then in one byte, the number of units supported is stored. This information is padded with seven zeroes.

Strategy routine is called whenever the device is needed and it is passed a request header. Request header stores the unit requested, the command code, space for return value and buffer pointers etc. Important command codes include 0 to initialize, 1 to check media, 2 to build a BIOS parameter block, 4 and 8 for read and write respectively. For every command the first 13 bytes of request header are same.

RH+0	BYTE	Length of request header
RH+1	BYTE	Unit requested
RH+2	BYTE	Command code
RH+3	BYTE	Driver's return code
RH+5	9 BYTES	Reserved

The request header details for different commands is listed below.

0 - Driver Initialization

Passed to driver

RH+18 DWORD Pointer to character after equal sign on CONFIG.SYS line that loaded driver (read-only)

```

RH+22  BYTE    Drive number for first unit of this block driver
          (0=A...)
Return from driver
RH+13  BYTE    Number of units (block devices only)
RH+14  DWORD   Address of first free memory above driver (break
          address)
RH+18  DWORD   BPB pointer array (block devices only)
1 - Media Check
RH+13  BYTE    Media descriptor byte
Return
RH+14  BYTE    Media change code
          -1 if disk changed
          0 if dont know whether disk changed
          1 if disk not changed
RH+15  DWORD   pointer to previous volume label if device attrib bit
          11=1 (open/close/removable media supported)
2 - Build BPB
RH+13  BYTE    Media descriptor byte
RH+14  DWORD   buffer address (one sector)
Return
RH+18  DWORD   pointer to new BPB
if bit 13 (ibm format) is set buffer is first sector of fat, otherwise
scrach space
4 - Read / 8 - Write / 9 - Write with verify
RH+13  BYTE    Media descriptor byte
RH+14  DWORD   transfer address
RH+18  WORD    byte or sector count
RH+20  WORD    starting sector number (for block devices)
Return
RH+18  WORD    actual byte or sectors transferred
RH+22  DWORD   pointer to volume label if error 0Fh is returned

```

The BIOS parameter block discussed above is a structure that provides parameters about the storage medium. It is stored in the first sector or the boot sector of the device. Its contents are listed below.

```

00-01  bytes per sector
02      sectors per allocation unit
03-04  Number of reserved sectors ( 0 based)
05      number of file allocation tables
06-07  max number of root directory entries
08-09  total number of sectors in medium
0A      media descriptor byte
0B-0C  number of sectors occupied by a single FAT
0D-0E  sectors per track (3.0 or later)
0F-10  number of heads (3.0 or later)
11-12  number of hidden sectors (3.0 or later)
13-14  high-order word of number of hidden sectors (4.0)
15-18  IF bytes 8-9 are zero, total number of sectors in medium
19-1E  Reserved should be zero

```

We will be building an example device driver that takes some RAM and expresses it as a secondary storage device to the operating system. Therefore a new drive is added and that can be browsed to, filed copied to and from just like ordinary drives expect that this drive is very fast as it is located in the RAM. This program cannot be directly executed since it is not a user program. This must be loaded by adding the line "device=filename.sys" in the "config.sys" file in the root directory.

Example 13.5

```

001 ; ram disk dos block device driver
002 header: dd -1 ; no next driver
003 dw 0x2000 ; driver attributes: block device
004 dw strategy ; offset of strategy routine
005 dw interrupt ; offset of interrupt routine
006 db 1 ; no of units supported
007 times 7 db 0 ; reserved
008
009 request: dd 0 ; space for request header
010
011 ramdisk: times 11 db 0 ; initial part of boot sector
012 bpb: dw 512 ; bytes per sector
013 db 1 ; sectors per cluster
014 dw 1 ; reserved sectors
015 db 1 ; fat copies
016 dw 48 ; root dir entries
017 dw 105 ; total sectors
018 db 0xfe, 0xff, 0xff ; media desc byte: fixed disk
019 dw 1 ; sectors per fat
020 times 482 db 0 ; remaining part of boot sector
021 db 0xfe, 0xff, 0xff ; special bytes at start of FAT
022 times 509 db 0 ; remaining FAT entries unused
023 times 103*512 db 0 ; 103 sectors for data
024 bpbptr: dw bpb ; array of bpb pointers
025
026 dispatch: dw init ; command 0: init
027 dw mediacheck ; command 1: media check
028 dw getbpb ; command 2: get bpb
029 dw unknown ; command 3: not handled
030 dw input ; command 4: input
031 dw unknown ; command 5: not handled
032 dw unknown ; command 6: not handled
033 dw unknown ; command 7: not handled
034 dw output ; command 8: output
035 dw output ; command 9: output with verify
036
037 ; device driver strategy routine
038 strategy: mov [cs:request], bx ; save request header offset
039 mov [cs:request+2], es ; save request header segment
040 retf
041
042 ; device driver interrupt routine
043 interrupt: push ax
044 push bx
045 push cx
046 push dx
047 push si
048 push di
049 push ds
050 push es
051
052 push cs
053 pop ds
054
055 les di, [request]
056 mov word [es:di+3], 0x0100
057 mov bl, [es:di+2]
058 mov bh, 0
059 cmp bx, 9
060 ja skip
061 shl bx, 1
062
063 call [dispatch+bx]
064
065 skip: pop es
066 pop ds
067 pop di
068 pop si
069 pop dx
070 pop cx
071 pop bx
072 pop ax
073 retf
074

```

```

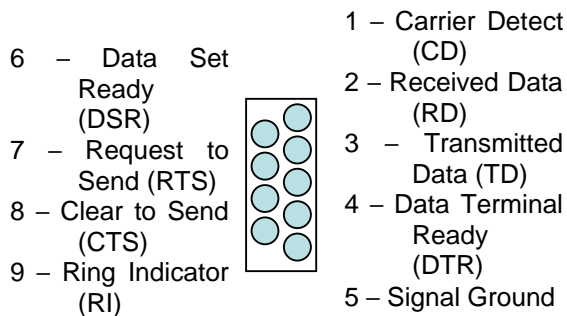
075 mediacheck: mov byte [es:di+14], 1
076               ret
077
078 getbpb:      mov word [es:di+18], bpb
079               mov [es:di+20], ds
080               ret
081
082 input:       mov ax, 512
083               mul word [es:di+18]
084               mov cx, ax
085
086               mov ax, 512
087               mul word [es:di+20]
088               mov si, ax
089               add si, ramdisk
090
091               les di, [es:di+14]
092               cld
093               rep movsb
094               ret
095
096 output:      mov ax, 512
097               mul word [es:di+18]
098               mov cx, ax
099
100               lds si, [es:di+14]
101               mov ax, 512
102               mul word [es:di+20]
103               mov di, ax
104               add di, ramdisk
105
106               push cs
107               pop es
108               cld
109               rep movsb
110 unknown:     ret
111
112 init:        mov ah, 9
113               mov dx, message
114               int 0x21
115
116               mov byte [es:di+13], 1
117               mov word [es:di+14], init
118               mov [es:di+16], ds
119               mov word [es:di+18], bpbptr
120               mov [es:di+20], ds
121               ret
122
123 message:     db 13, 10, 'RAM Disk Driver loaded', 13, 10, '$'

```


Serial Port Programming

7.1. INTRODUCTION

Serial port is a way of communication among two devices just like the parallel port. The basic difference is that whole bytes are sent from one place to another in case of parallel port while the bits are sent one by one on the serial port in a specially formatted fashion. The serial port connection is a 9pin DB-9 connector with pins assigned as shown below.



We have made a wire that connects signal ground of the two connectors, the TD of one to the RD of the other and the RD of one to the TD of the other. This three wire connection is sufficient for full duplex serial communication. The data on the serial port is sent in a standard format called RS232 communication. The data starts with a 1 bit called the start bit, then five to eight data bits, an optional parity bit, and one to two 0 bits called stop bits. The number of data bits, parity bits, and the number of stop bits have to be configured at both ends. Also the duration of a bit must be precisely known at both ends called the baud rate of the communication.

The BIOS INT 14 provides serial port services. We will use a mix of BIOS services and direct port access for our example. A major limitation in using BIOS is that it does not allow interrupt driven data transfer, i.e. we are interrupted whenever a byte is ready to be read or a byte can be transferred since the previous transmission has completed. To achieve this we have to resort to direct port access. Important BIOS services regarding the serial port are discussed below.


```

INT 14 - SERIAL - INITIALIZE PORT
AH = 00h
AL = port parameters
DX = port number (00h-03h)
Return:
AH = line status
AL = modem status

```

Every bit of line status conveys different information. From most significant to least significant, the meanings are timeout, transmitter shift register empty, transmitter holding register empty, break detect, receiver ready, overrun, parity error, and framing error. Modem status is not used in direct serial communication. The port parameters in AL consist of the baud

rate, parity scheme, number of stop bits, and number of data bits. The description of various bits is as under.



INT 14 - SERIAL - WRITE CHARACTER TO PORT

AH = 01h
 AL = character to write
 DX = port number (00h-03h)
 Return:
 AH bit 7 = error flag
 AH bits 6-0 = port status

INT 14 - SERIAL - READ CHARACTER FROM PORT

AH = 02h
 DX = port number (00h-03h)
 Return:
 AH = line status
 AL = received character if AH bit 7 clear

INT 14 - SERIAL - GET PORT STATUS

AH = 03h
 DX = port number (00h-03h)
 Return:
 AH = line status
 AL = modem status

Serial port is also accessible via I/O ports. COM1 is accessible via ports 3F8-3FF while COM2 is accessible via 2F8-2FF. The first register at 3F8 (or 2F8 for the other port) is the transmitter holding register if written to and the receiver buffer register if read from. Other registers of our interest include 3F9 whose bit 0 must be set to enable received data available interrupt and bit 1 must be set to enable transmitter holding register empty interrupt. Bit 0 of 3FA is set if an interrupt is pending and its bits 1-3 identify the cause of the interrupt. The three bit causes are as follows.

110	(16550, 82510) timeout interrupt pending
101	(82510) timer interrupt
100	(82510) transmit machine
011	receiver line status interrupt. priority=highest
010	received data available register interrupt. priority=second
001	transmitter holding register empty interrupt. priority=third
000	modem status interrupt. priority=fourth

The register at 3FB is line control register while the one at 3FD is line status register. The line status register has the same bits as returned in line status by the get port status BIOS interrupt however the most significant bit

is reserved in this case instead of signaling a timeout. The register at 3FC is the modem control register. Bit 3 of this register must be set to enable interrupt generation by the serial port.

7.2. SERIAL COMMUNICATION

We give an example where two computers are connected using a serial cable made just as described above. The program is to be run on both computers. After that whatever is typed on one computer appears on the screen of the other.

Example 14.1

```

001 ; a program using serial port to transfer data back and forth
002 [org 0x0100]
003         jmp  start
004
005 screenpos:  dw  0           ; where to display next character
006
007 ; subroutine to clear the screen
008 clrscr:    push es
009            push ax
010            push cx
011            push di
012
013            mov ax, 0xb800
014            mov es, ax       ; point es to video base
015            xor di, di       ; point di to top left column
016            mov ax, 0x0720   ; space char in normal attribute
017            mov cx, 2000     ; number of screen locations
018
019            cld               ; auto increment mode
020            rep stosw        ; clear the whole screen
021
022            pop di
023            pop cx
024            pop ax
025            pop es
026            ret
027
028 serial:    push ax
029            push bx
030            push dx
031            push es
032
033            mov dx, 0x3FA     ; interrupt identification register
034            in  al, dx        ; read register
035            and al, 0x0F      ; leave lowernibble only
036            cmp al, 4         ; is receiver data available
037            jne skipall      ; no, leave interrupt handler
038
039            mov dx, 0x3F8     ; data register
040            in  al, dx        ; read character
041
042            mov dx, 0xB800
043            mov es, dx        ; point es to video memory
044            mov bx, [cs:screenpos] ; get current screen position
045            mov [es:bx], al    ; write character on screen
046            add word [cs:screenpos], 2 ; update screen position
047            cmp word [cs:screenpos], 4000 ; is the screen full
048            jne skipall      ; no, leave interrupt handler
049
050            call clrscr       ; clear the screen
051            mov word [cs:screenpos], 0 ; reset screen position
052
053 skipall:    mov al, 0x20
054            out 0x20, al      ; end of interrupt
055
056            pop es
057            pop dx
058            pop bx
059            pop ax
060            iret

```

```

061
062      start:      call clrscr          ; clear the screen
063
064                mov  ah, 0             ; initialize port service
065                mov  al, 0xE3          ; line settings = 9600, 8, N, 1
066                xor  dx, dx            ; port = COM1
067                int  0x14              ; BIOS serial port services
068
069                xor  ax, ax
070                mov  es, ax            ; point es to IVT base
071                mov  word [es:0x0C*4], serial
072                mov  [es:0x0C*4+2], cs ; hook serial port interrupt
073
074                mov  dx, 0x3FC         ; modem control register
075                in   al, dx            ; read register
076                or   al, 8             ; enable bit 3 (OUT2)
077                out  dx, al            ; write back to register
078
079                mov  dx, 0x3F9         ; interrupt enable register
080                in   al, dx            ; read register
081                or   al, 1             ; receiver data interrupt enable
082                out  dx, al            ; write back to register
083
084                in   al, 0x21          ; read interrupt mask register
085                and  al, 0xEF          ; enable IRQ 4
086                out  0x21, al         ; write back to register
087
088      main:      mov  ah, 0             ; read key service
089                int  0x16              ; BIOS keyboard services
090                push ax                ; save key for later use
091
092      retest:    mov  ah, 3             ; get line status
093                xor  dx, dx            ; port = COM1
094                int  0x14              ; BIOS keyboard services
095                and  ah, 32            ; trasmitter holding register empty
096                jz   retest           ; no, test again
097
098                pop  ax                ; load saved key
099                mov  dx, 0x3F8         ; data port
100                out  dx, al            ; send on serial port
101
102                jmp  main

```

8

Protected Mode Programming

8.1. INTRODUCTION

Till now we have been discussing the 8088 architecture which was a 16bit processor. Newer processors of the Intel series provide 32bit architecture. Till now we were in real mode of a newer processor which is basically a compatibility mode making the newer processor just a faster version of the original 8088. Switching processor in the newer 32bit mode is a very easy task. Just turn on the least significant bit of a new register called CR0 (Control Register 0) and the processor switches into 32bit mode called protected mode. However manipulations in the protected mode are very different from those in the real mode.

All registers in 386 have been extended to 32bits. The new names are EAX, EBX, ECX, EDX, ESI, EDI, ESP, EBP, EIP, and EFLAGS. The original names refer to the lower 16bits of these registers. A 32bit address register can access upto 4GB of memory so memory access has increased a lot.

As regards segment registers the scheme is not so simple. First of all we call them segment selectors instead of segment registers and they are still 16bits wide. We are also given two other segment selectors FS and GS for no specific purpose just like ES.

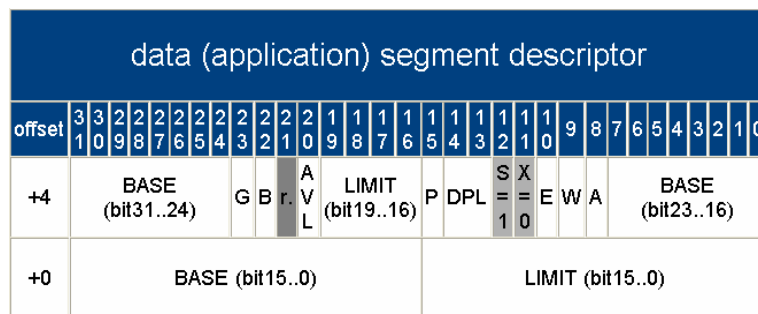
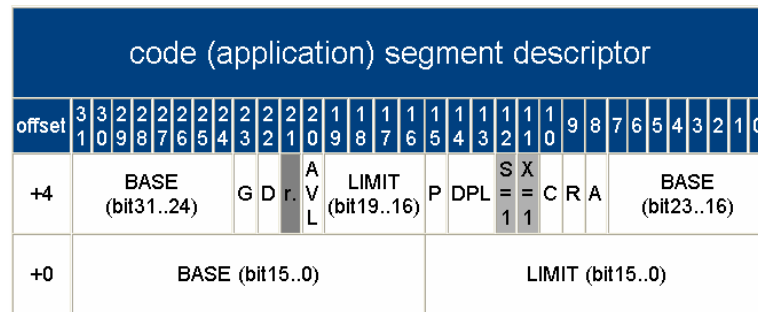
The working of segment registers as being multiplied by 10 and added into the offset for obtaining the physical address is totally changed. Now the selector is just an index into an array of segment descriptors where each descriptor describes the base, limit, and attributes of a segment. Role of selector is to select on descriptor from the table of descriptors and the role of descriptor is to define the actual base address. This decouples the selection and actual definition which is needed in certain protection mechanisms introduced into this processor. For example an operating system can define the possible descriptors for a program and the program is bound to select one of them and nothing else. This sentence also hints that the processor has some sense of programs that can or cannot do certain things like change this table of descriptors. This is called the privilege level of the program and varies for 0 (highest privilege) to 3 (lowest privilege). The format of a selector is shown below.



The table index (TI) is set to 0 to access the global table of descriptors called the GDT (Global Descriptor Table). It is set to 1 to access another table, the local descriptor table (LDT) that we will not be using. RPL is the requested privilege level that ranges from 0-3 and informs what privilege level

the program wants when using this descriptor. The 13bit index is the actual index into the GDT to select the appropriate descriptor. 13 bits mean that a maximum of 8192 descriptors are possible in the GDT.

The GDT itself is an array of descriptors where each descriptor is an 8byte entry. The base and limit of GDT is stored in a 48bit register called the GDTR. This register is loaded with a special instruction LGDT and is given a memory address from where the 48bits are fetched. The first entry of the GDT must always be zero. It is called the null descriptor. After that any number of entries upto a maximum of 8191 can follow. The format of a code and data descriptor is shown below.



The 32bit base in both descriptors is scattered into different places because of compatibility reasons. The limit is stored in 20 bits but the G bit defines that the limit is in terms of bytes of 4K pages therefore a maximum of 4GB size is possible. The P bit must be set to signal that this segment is present in memory. DPL is the descriptor privilege level again related to the protection levels in 386. D bit defines that this segment is to execute code is 16bit mode or 32bit mode. C is conforming bit that we will not be using. R signals that the segment is readable. A bit is automatically set whenever the segment is accessed. The combination of S (system) and X (executable) tell that the descriptors are a code or a data descriptor. B (big) bit tells that if this data segment is used as stack SP is used or ESP is used.

Our first example is a very rudimentary one that just goes into protected mode and prints an A on the screen by directly accessing 000B8000.

Example 15.1	
001	[org 0x0100]
002	jmp start
003	
004	gdt: dd 0x00000000, 0x00000000 ; null descriptor
005	dd 0x0000FFFF, 0x00CF9A00 ; 32bit code
006	; \--/\--/ \/ \\
007	; +--- Base (16..23)=0 fill later
008	; +--- X=1 C=0 R=1 A=0
009	; +--- P=1 DPL=0 S=1
010	; +--- Limit (16..19) = F
011	; +--- G=1 D=1 r=0 AVL=0

```

012      ;      |      |      +--- Base (24..31) = 0
013      ;      |      |      +--- Limit (0..15) = FFFF
014      ;      |      |      +--- Base (0..15)=0 fill later
015      dd      0x0000FFFF, 0x00CF9200      ; data
016      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
017      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
018      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
019      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
020      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
021      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
022      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
023      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
024      ;      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
025
026      gdtreg:      dw      0x17      ; 16bit limit
027                  dd      0      ; 32bit base (filled later)
028
029      stack:      times 256 dd 0 ; for use in p-mode
030      stacktop:
031
032      start:      mov      ax, 0x2401
033                  int      0x15      ; enable A20
034
035                  xor      eax, eax
036                  mov      ax, cs
037                  shl      eax, 4
038                  mov      [gdt+0x08+2], ax
039                  shr      eax, 16
040                  mov      [gdt+0x08+4], al      ; fill base of code desc
041
042                  xor      edx, edx
043                  mov      dx, cs
044                  shl      edx, 4
045                  add      edx, stacktop      ; edx = stack top for p-
046      mode
047
048                  xor      eax, eax
049                  mov      ax, cs
050                  shl      eax, 4
051                  add      eax, gdt
052                  mov      [gdtreg+2], eax      ; fill phy base of gdt
053                  lgdt      [gdtreg]      ; load gdtr
054
055                  mov      eax, cr0
056                  or      eax, 1
057
058                  cli      ; MUST disable interrupts
059                  mov      cr0, eax      ; P-MODE ON
060                  jmp      0x08:pstart      ; load cs
061
062      ;;;; 32bit protected mode ;;;;
063
064      [bits 32] ; ask assembler to generate 32bit code
065      pstart:      mov      eax, 0x10
066                  mov      ds, ax
067                  mov      es, ax      ; load other seg regs
068                  mov      fs, ax      ; flat memory model
069                  mov      gs, ax
070                  mov      ss, ax
071                  mov      esp, edx
072
073                  mov      byte [0x000b8000], 'A' ; direct poke at video
074                  jmp      $      ; hang around
075

```

Gate A20 is a workaround for a bug that is not detailed here. The BIOS call will simply enable it to open the whole memory for us. Another important thing is that the far jump we used loaded 8 into CS but CS is now a selector so it means Index=1, TI=0, and RPL=0 and therefore the actual descriptor loaded is the one at index 1 in the GDT.

8.2. 32BIT PROGRAMMING

Our next example is to give a flavour of 32bit programming. We have written the printstr function for real and for protected mode. The availability of larger registers and flexible addressing rules allows writing a much comprehensive version of the code. Also offsets to parameters and default widths change.

Example 15.2	
001	[org 0x0100]
002	jmp start
003	
004	gdt: dd 0x00000000, 0x00000000 ; null descriptor
005	dd 0x0000FFFF, 0x00CF9A00 ; 32bit code
006	dd 0x0000FFFF, 0x00CF9200 ; data
007	
008	gdtreg: dw 0x17 ; 16bit limit
009	dd 0 ; 32bit base
010	
011	rstring: db 'In Real Mode...', 0
012	pstring: db 'In Protected Mode...', 0
013	
014	stack: times 256 dd 0 ; 1K stack
015	stacktop:
016	
017	printstr: push bp ; real mode print string
018	mov bp, sp
019	push ax
020	push cx
021	push si
022	push di
023	push es
024	
025	mov di,[bp+4] ;load string address
026	mov cx,0xffff ;load maximum possible size in cx
027	xor al,al ;clear al reg
028	repne scasb ;repeat scan
029	mov ax,0xffff ;
030	sub ax,cx ;calculate length
031	dec ax ;off by one, as it includes zero
032	mov cx,ax ;move length to counter
033	
034	mov ax, 0xb800
035	mov es, ax ; point es to video base
036	mov ax,80 ;its a word move, clears ah
037	mul byte [bp+8] ;its a byte mul to calc y offset
038	add ax,[bp+10] ;add x offset
039	shl ax,1 ;mul by 2 to get word offset
040	mov di,ax ;load pointer
041	
042	mov si, [bp+4] ; string to be printed
043	mov ah, [bp+6] ; load attribute
044	
045	cld ; set auto increment mode
046	nextchar: lodsb ;load next char and inc si by 1
047	stosw ;store ax and inc di by 2
048	loop nextchar
049	
050	pop es
051	pop di
052	pop si
053	pop cx
054	pop ax
055	pop bp
056	ret 8
057	
058	start: push byte 0 ; 386 can directly push
059	immediates
060	push byte 10
061	push byte 7
062	push word rstring
063	call printstr
064	


```

065      mov ax, 0x2401
066      int 0x15          ; enable a20
067
068      xor eax, eax
069      mov ax, cs
070      shl eax, 4
071      mov [gdt+0x08+2], ax
072      shr eax, 16
073      mov [gdt+0x08+4], al      ; set base of code desc
074
075      xor edx, edx
076      mov dx, cs
077      shl edx, 4
078      add edx, stacktop      ; stacktop to be used in p-mode
079
080      xor ebx, ebx
081      mov bx, cs
082      shl ebx, 4
083      add ebx, pstring      ; pstring to be used in p-mode
084
085      xor eax, eax
086      mov ax, cs
087      shl eax, 4
088      add eax, gdt
089      mov [gdtreg+2], eax      ; set base of gdt
090      lgdt [gdtreg]          ; load gdt
091
092      mov eax, cr0
093      or  eax, 1
094
095      cli                    ; disable interrupts
096      mov cr0, eax          ; enable protected mode
097      jmp 0x08:pstart      ; load cs
098
099      ;;;; 32bit protected mode ;;;;
100
101      [bits 32]
102      pprintstr:  push ebp          ; p-mode print string routine
103                  mov ebp, esp
104                  push eax
105                  push ecx
106                  push esi
107                  push edi
108
109                  mov edi, [ebp+8] ;load string address
110                  mov ecx, 0xffffffff ;load maximum possible size in cx
111                  xor al, al      ;clear al reg
112                  repne scasb     ;repeat scan
113                  mov eax, 0xffffffff ;
114                  sub eax, ecx     ;calculate length
115                  dec eax         ;off by one, as it includes zero
116                  mov ecx, eax     ;move length to counter
117
118                  mov eax, 80      ;its a word move, clears ah
119                  mul byte [ebp+16] ;its a byte mul to calc y
120
121      offset      add eax, [ebp+20] ;add x offset
122                  shl eax, 1       ;mul by 2 to get word offset
123                  add eax, 0xb8000
124                  mov edi, eax     ;load pointer
125
126
127                  mov esi, [ebp+8] ; string to be printed
128                  mov ah, [ebp+12] ; load attribute
129
130      pnextchar:  cld                ; set auto increment mode
131                  lodsb             ;load next char and inc si by 1
132                  stosw            ;store ax and inc di by 2
133                  loop pnextchar
134
135                  pop edi
136                  pop esi
137                  pop ecx
138                  pop eax
139                  pop ebp
140                  ret 16           ; 4 args now mean 16 bytes
141

```

```

142      pstart:      mov     ax, 0x10          ; load all seg regs to 0x10
143                  mov     ds, ax            ; flat memory model
144                  mov     es, ax
145                  mov     fs, ax
146                  mov     gs, ax
147                  mov     ss, ax
148                  mov     esp, edx          ; load saved esp on stack
149
150                  push     byte 0
151                  push     byte 11
152                  push     byte 7
153                  push     ebx
154                  call     pprintstr        ; call p-mode print string
155
156      routine
157
158                  mov     eax, 0x000b8000
159                  mov     ebx, '/-\\|'
160
161      nextsymbol:   mov     [eax], bl
162                  mov     ecx, 0x00FFFFFF
163                  loop    $
164                  ror     ebx, 8
165                  jmp     nextsymbol

```

8.3. VESA LINEAR FRAME BUFFER

As an example of accessing a really large area of memory for which protected mode is a necessity, we will be accessing the video memory in high resolution and high color graphics mode where the necessary video memory is alone above a megabyte. We will be using the VESA VBE 2.0 for a standard for these high resolution modes.

VESA is the Video Electronics Standards Association and VBE is the set of Video BIOS Extensions proposed by them. The VESA VBE 2.0 standard includes a linear frame buffer mode that we will be using. This mode allows direct access to the whole video memory. Some important VESA services are listed below.

```

INT 10 - VESA - Get SuperVGA Information
AX = 4F00h
ES:DI -> buffer for SuperVGA information
Return:
AL = 4Fh if function supported
AH = status

INT 10 - VESA - Get SuperVGA Mode Information
AX = 4F01h
CX = SuperVGA video mode
ES:DI -> 256-byte buffer for mode information
Return:
AL = 4Fh if function supported
AH = status
ES:DI filled if no error

INT 10 - VESA - Set VESA Video Mode
AX = 4F02h
BX = new video mode
Return:
AL = 4Fh if function supported
AH = status

```

One of the VESA defined modes is 8117 which is a 1024x768 mode with 16bit color and a linear frame buffer. The 16 color bits for every pixel are organized in 5:6:5 format with 5 bits for red, 6 for green, and 5 for blue. This makes 32 shades of red and blue and 64 shades of green and 64K total

possible colors. The 32bit linear frame buffer base address is available at offset 28 in the mode information buffer. Our example will produces shades of green on the screen and clear them and again print them in an infinite loop with delays inbetween.

Example 15.3

```

001  [org 0x0100]
002      jmp  start
003
004  modeblock:  times 256 db 0
005
006  gdt:      dd  0x00000000, 0x00000000  ; null descriptor
007           dd  0x0000FFFF, 0x00CF9A00  ; 32bit code
008           dd  0x0000FFFF, 0x00CF9200  ; data
009
010  gdtreg:   dw  0x17                      ; 16bit limit
011           dd  0                        ; 32bit base
012
013  stack:    times 256 dd 0                ; 1K stack
014  stacktop:
015
016  start:    mov  ax, 0x4f01                ; get vesa mode information
017           mov  cx, 0x8117                ; 1024*768*64K linear frame
018  buffer
019           mov  di, modeblock
020           int  0x10
021           mov  esi, [modeblock+0x28]     ; save frame buffer base
022
023           mov  ax, 0x4f02                ; set vesa mode
024           mov  bx, 0x8117
025           int  0x10
026
027           mov  ax, 0x2401
028           int  0x15                      ; enable a20
029
030           xor  eax, eax
031           mov  ax, cs
032           shl  eax, 4
033           mov  [gdt+0x08+2], ax
034           shr  eax, 16
035           mov  [gdt+0x08+4], al          ; set base of code desc
036
037           xor  edx, edx
038           mov  dx, cs
039           shl  edx, 4
040           add  edx, stacktop             ; stacktop to be used in p-mode
041
042           xor  eax, eax
043           mov  ax, cs
044           shl  eax, 4
045           add  eax, gdt
046           mov  [gdtreg+2], eax           ; set base of gdt
047           lgdt [gdtreg]                 ; load gdt
048
049           mov  eax, cr0
050           or   eax, 1
051
052           cli                                ; disable interrupts
053           mov  cr0, eax                  ; enable protected mode
054           jmp  0x08:pstart              ; load cs
055
056  ;;;; 32bit protected mode ;;;;
057
058  [bits 32]
059  pstart:   mov  ax, 0x10                ; load all seg regs to 0x10
060           mov  ds, ax                  ; flat memory model
061           mov  es, ax
062           mov  fs, ax
063           mov  gs, ax
064           mov  ss, ax
065           mov  esp, edx                 ; load saved esp on stack
066
067  ll:      xor  eax, eax

```

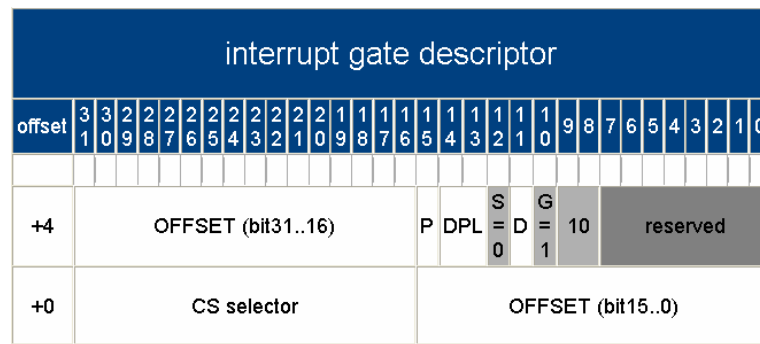
```

068      mov     edi, esi
069      mov     ecx, 1024*768*2/4      ; divide by 4 as dwords
070      cld
071      rep     stosd
072
073      mov     eax, 0x07FF07FF
074      mov     ecx, 32                ; no of bands
075      mov     edi, esi
076
077      12:      push    ecx
078      mov     ecx, 768*16            ; band width = 32
079      lines
080      cld
081      rep     stosd
082
083      mov     ecx, 0x000FFFFF        ; small wait
084      loop    $
085      pop     ecx
086
087      sub     eax, 0x00410041
088      loop    12
089
090      mov     ecx, 0xFFFFFFFF        ; long wait
091      loop    $
092      jmp     11
093

```

8.4. INTERRUPT HANDLING

Handling interrupts in protected mode is also different. Instead of the IVT at physical address 0 there is the IDT (interrupt descriptor table) located at physical address stored in IDTR, a special purpose register. The IDTR is also a 48bit register similar in structure to the GDTR and loaded with another special instruction LGDT. The format of the interrupt descriptor is as shown below.



The P and DPL have the same meaning as in data and code descriptors. The S bit tells that this is a system descriptor while the 1110 following it tells that it is a 386 interrupt gate. Our example hooks the keyboard and timer interrupts and displays certain things on the screen to show that they are working.

Example 15.4	
001	[org 0x0100]
002	jmp start
003	
004	gdt: dd 0x00000000, 0x00000000 ; null descriptor
005	dd 0x0000FFFF, 0x00CF9A00 ; 32bit code
006	dd 0x0000FFFF, 0x00CF9200 ; data
007	
008	gdtreg: dw 0x17 ; 16bit limit
009	dd 0 ; 32bit base

```

010
011 idt:      times 8 dw 0x0008, 0x8e00, 0x0000
012          dw  timer, 0x0008, 0x8e00, 0x0000
013          \---/ \---/ ||\ \---/
014          |      |    ||| +----- offset bits 16..32
015          |      |    ||| +----- reserved
016          |      |    ||| +----- Type=E 386 Interrupt Gate
017          |      |    ||| +---- P=1 DPL=00 S=0
018          |      |    ||| +----- selector
019          |      |    ||| +----- offset bits 0..15
020          dw  keyboard, 0x0008, 0x8e00, 0x0000
021          times 246 dw 0x0008, 0x8e00, 0x0000
022
023 idtreg:   dw  0x07FF
024          dd  0
025
026 stack:    times 256 dd 0          ; 1K stack
027 stacktop:
028
029 start:    mov  ax, 0x2401
030          int  0x15          ; enable a20
031
032          xor  eax, eax
033          mov  ax, cs
034          shl  eax, 4
035          mov  [gdt+0x08+2], ax
036          shr  eax, 16
037          mov  [gdt+0x08+4], al    ; set base of code desc
038
039          xor  edx, edx
040          mov  dx, cs
041          shl  edx, 4
042          add  edx, stacktop      ; stacktop to be used in p-mode
043
044          xor  eax, eax
045          mov  ax, cs
046          shl  eax, 4
047          add  eax, gdt
048          mov  [gdtreg+2], eax    ; set base of gdt
049          lgdt [gdtreg]          ; load gdtr
050
051          xor  eax, eax
052          mov  ax, cs
053          shl  eax, 4
054          add  eax, idt
055          mov  [idtreg+2], eax    ; set base of idt
056
057          cli          ; disable interrupts
058          lidt [idtreg]        ; load idtr
059
060          mov  eax, cr0
061          or   eax, 1
062          mov  cr0, eax        ; enable protected mode
063
064          jmp  0x08:pstart      ; load cs
065
066      ;;;; 32bit protected mode ;;;;
067
068 [bits 32]
069 unhandled:  iret
070
071 timer:      push  eax
072
073             inc  byte [0x000b8000]
074
075             mov  al, 0x20
076             out  0x20, al
077             pop  eax
078             iret
079
080 keyboard:   push  eax
081
082             in   al, 0x60
083             mov  ah, al
084             and  al, 0x0F
085             shr  ah, 4

```

```

086      add ax, 0x3030
087      cmp al, 0x39
088      jbe skip1
089      add al, 7
090      skip1: cmp ah, 0x39
091      jbe skip2
092      add ah, 7
093      skip2: mov [0x000b809C], ah
094      mov [0x000b809E], al
095
096      skipkb: mov al, 0x20
097      out 0x20, al
098      pop eax
099      iret
100
101      pstart: mov ax, 0x10      ; load all seg regs to 0x10
102      mov ds, ax      ; flat memory model
103      mov es, ax
104      mov fs, ax
105      mov gs, ax
106      mov ss, ax
107      mov esp, edx      ; load saved esp on stack
108
109      mov al, 0xFC
110      out 0x21, al      ; no unexpected int comes
111
112      sti      ; interrupts are okay now
113
114      jmp $

```

EXERCISES

- Write very brief and to-the-point answers.
 - Why loading idtr with a value appropriate for real mode is necessary while gdtr is not?
 - What should we do in protected mode so that when we turn protection off, we are in unreal mode?
 - If the line `jmp code:next` is replaced with `call code:somefun`, the prefetch queue is still emptied. What problem will occur when `somefun` will return?
 - How much is ESP decremented when an interrupt arrives. This depends on whether we are in 16-bit mode or 32-bit. Does it depend on any other thing as well? If yes, what?
 - Give two instructions that change the TR register.
- Name the following descriptors like code descriptor, data descriptor, interrupt gate etc.

```

gdt:  dd 0x00000000, 0x00000000
      dd 0x00000000, 0x00000000
      dd 0x80000fa0, 0x0000820b
      dd 0x0000ffff, 0x00409a00
      dd 0x80000000, 0x0001d20b

```

- Using the above GDT, which of the following values, when moved into DS will cause an exception and why.

```

0x00
0x08
0x10
0x18
0x28
0x23

```

- Using the above GDT, if DS contains 0x20, which of the following offsets will cause an exception on read access?

```

0x0ffff
0x10000
0x10001

```

5. The following function is written in 32-bit code for a 16-bit stack. Against every instruction, write the prefixes generated before that instruction. Prefixes can be address size, operand size, repeat, or segment override. Then rewrite the code such that no prefixes are generated considering that this is assembled and executed in 32-bit mode. Don't care for retaining register values. The function copies specified number of DWORDs between two segments.

```
[bits 32]
memcpy:    mov     bp, sp
           lds     esi, [bp+4]          ; source address
           les     edi, [bp+10]         ; destination address
           mov     cx, [bp+16]          ; count of DWORDs to move
           shl     cx, 1                ; make into count of WORDs
L1:         mov     dx, [si]
           mov     [es:di], dx
           dec     cx
           jnz     L1
           ret
```

6. Rewrite the following scheduler so that it schedules processes stored in readyQ, where enqueue and deque functions are redefined and readyQ contains TSS selectors of processes to be multitasked. Remember you can't use a register as a segment in a jump (eg `jmp ax:0`) but you can jump to an indirect address (eg `jmp far [eax]`) where `eax` points to a six-byte address. Declare any variables you need.

```
scheduler: mov     al, 0x20
           jmp     USERONESEL:0
           out     0x20, al
           mov     byte [USERONEDESC+5], 0x89
           jmp     USERTWOSEL:0
           out     0x20, al
           mov     byte [USERTWODESC+5], 0x89
           jmp     scheduler
```

7. Protected mode has specialized mechanism for multitasking using task state segments but the method used in real mode i.e. saving all registers in a PCB, selecting the next PCB and loading all registers from there is still applicable. Multitask two tasks in protected mode multitasking without TSS. Assume that all processes are at level zero so no protection issues arise. Be careful to save the complete state of the process.
8. Write the following descriptors.
- 32 bit, conforming, execute-only code segment at level 2, with base at 6MB and a size of 4MB.
 - 16 bit, non-conforming, readable code segment at level 0, with base at 1MB and a size of 10 bytes.
 - Read only data segment at level 3, with base at 0 and size of 1MB.
 - Interrupt Gate with selector 180h and offset 11223344h.
9. Write physical addresses for the following accesses where CS points to the first descriptor above, DS to the second, ES to the third, EBX contains 00010000h, and ESI contains 00020000h
- [bx+si]
 - [ebx+esi-2ffffh]
 - [es:ebx-10h]
10. Which of the following will cause exceptions and why. The registers have the same values as the last question.
- `mov eax, [cs:10000h]`
 - `mov [es:esi:100h], ebx`
 - `mov ax, [es:ebx]`
11. Give short answers.
- How can a GPF (General protection fault) occur while running the following code

```
push es
pop  es
```

- b. How can a GPF occur during the following instruction? Give any two reasons.
`jmp 10h:100h`
 - c. What will happen if we call interrupt 80h after loading out IDT and before switching to protected mode?
 - d. What will happen if we call interrupt 80h after switching into protected mode but before making a far jump?
12. Write the following descriptors. Assume values for attributes not specifically mentioned.
- a. Write a 32-bit data segment with 1 GB base and 1 GB limit and a privilege level of 2.
 - b. Readable 16-bit code descriptor with 1 MB base and 1 MB limit and a privilege level of 1.
 - c. Interrupt gate given that the handler is at 48h:12345678h and a privilege level of 0.
13. Describe the following descriptors. Give their type and the value of all their fields.
- ```
dd 01234567h, 789abcdeh
dd 30405060h, 70809010h
dd 00aabb00h, 00ffee00h
```
14. Make an EXE file, switch into protected mode, rotate an asterisk on the border of the screen, and return to real mode when the border is traversed.



# 9

# Interfacing with High Level Languages

## 9.1. CALLING CONVENTIONS

---

To interface an assembly routine with a high level language program means to be able to call functions back and forth. And to be able to do so requires knowledge of certain behavior of the HLL when calling functions. This behavior of calling functions is called the calling conventions of the language. Two prevalent calling conventions are the C calling convention and the Pascal calling convention.

### What is the naming convention

C prepends an underscore to every function or variable name while Pascal translates the name to all uppercase. C++ has a weird name mangling scheme that is compiler dependent. To avoid it C++ can be forced to use C style naming with extern "C" directive.

### How are parameters passed to the routine

In C parameters are pushed in reverse order with the rightmost being pushed first. While in Pascal they are pushed in proper order with the leftmost being pushed first.

### Which registers must be preserved

Both standards preserve EBX, ESI, EDI, EBP, ESP, DS, ES, and SS.

### Which registers are used as scratch

Both standards do not preserve or guarantee the value of EAX, ECX, EDX, FS, GS, EFLAGS, and any other registers.

### Which register holds the return value

Both C and Pascal return upto 32bit large values in EAX and upto 64bit large values in EDX:EAX.

### Who is responsible for removing the parameters

In C the caller removes the parameter while in Pascal the callee removes them. The C scheme has reasons pertaining to its provision for variable number of arguments.

## 9.2. CALLING C FROM ASSEMBLY

---

For example we take a function divide declared in C as follows.

```
int divide(int dividend, int divisor);
```

To call this function from assembly we have to write.

```
push dword [mydivisor]
```

```
push dword [mydividend]
call _divide
add esp, 8
; EAX holds the answer
```

Observe the order of parameters according to the C calling conventions and observe that the caller cleared the stack. Now take another example of a function written in C as follows.

```
void swap(int* p1, int* p2)
{
 int temp = *p1;
 *p1 = *p2;
 *p2 = temp;
}
```

To call it from assembly we have to write this.

```
[section .text]
extern _swap
x: dd 4
y: dd 7

push dword y
push dword x
call _swap ; will only retain the specified registers
add esp, 8
```

Observe how pointers were initialized appropriately. The above function swap was converted into assembly by the gcc compiler as follows.

```
; swap generated by gcc with no optimizations (converted to Intel
syntax)
; 15 instructions AND 13 memory accesses
_swap:
 push ebp
 mov ebp, esp
 sub esp, 4 ; space created for temp

 mov eax, [ebp+8]
 mov eax, [eax]
 mov [ebp-4], eax ; temp = *p1

 mov edx, [ebp+8]
 mov eax, [ebp+12]
 mov eax, [eax]
 mov [edx], eax ; *p1 = *p2

 mov edx, [ebp+12]
 mov eax, [ebp-4]
 mov [edx], eax ; *p2 = temp

 leave ;;;; EQUIVALENT TO mov esp, ebp AND pop ebp ;;;;
 ret
```

If we turn on optimizations the same function is compiled into the following code.

```

; generated with full optimization by gcc compiler
; 12 instructions AND 11 memory accesses
_swap:
 push ebp
 mov ebp, esp
 push ebx

 mov edx, [ebp+8]
 mov ecx, [ebp+12]
 mov ebx, [edx]
 mov eax, [ecx]
 mov [edx], eax
 mov [ecx], ebx

 pop ebx
 pop ebp
 ret

```

### 9.3. CALLING ASSEMBLY FROM C

We now write a hand optimized version in assembly. Our version is only 6 instructions and 6 memory accesses.

#### Example 16.1

```

001 [section .text]
002 global _swap
003 _swap: mov ecx,[esp+4] ; copy parameter p1 to ecx
004 mov edx,[esp+8] ; copy parameter p2 to edx
005 mov eax,[ecx] ; copy *p1 into eax
006 xchg eax,[edx] ; exchange eax with *p2
007 mov [ecx],eax ; copy eax into *p1
008 ret ; return from this function

```

We assemble the above program with the following command.

- `nasm -f win32 swap.asm`

This produces a `swap.obj` file. The format directive told the assembler that it is to be linked with a 32bit Windows executable. The linking process involves resolving imported symbols of one object files with export symbols of another. In NASM an imported symbol is declared with the `extern` directive while and exported symbol is declared with the `global` directive.

We write the following program in C to call this assembly routine. We should have provided the `swap.obj` file to the C linker otherwise an unresolved external symbol error will come.

#### Example 16.1

```

001 #include <stdio.h>
002
003 void swap(int* p1, int* p2);
004
005 int main()
006 {
007 int a = 10, b = 20;
008 printf("a=%d b=%d\n", a, b);
009 swap(&a, &b);
010 printf("a=%d b=%d\n", a, b);
011 system("PAUSE");
012 return 0;
013 }

```

**EXERCISES**

---

1. Write a traverse function in assembly, which takes an array, the number of elements in the array and the address of another function to be called for each member of the array. Call the function from a C program.
2. Make the linked list functions made in Exercise 5.XX available to C programs using the following declarations.

```
struct node {
 int data;
 struct node* next;
};
void init(void);
struct node* createlist(void);
void insertafter(struct node*, int);
void deleteafter(struct node*);
void deletelist(struct node*);
```

3. Add two functions to the above program implemented in C. The function “printnode” should print the data in the passed node using printf, while “countfree” should count the number of free nodes by traversing the free list starting from the node address stored in firstfree.

```
void printnode(struct node*);
void countfree(void);
```

4. Add the function “printlist” to the above program and implement in assembly. This function should traverse the list whose head is passed as parameter and for each node containing data (head is dummy and doesn’t contain data) calls the C function printnode to actually print the contained data.

```
void printlist(struct node*);
```

5. Modify the createlist and deletelist functions in the above program to increment and decrement an integer variable “listcount” declared in C to maintain a count of linked lists present.

# 10

## Comparison with Other Processors

We emphasized that assembly language has to be learned once and every processor can be programmed by that person. To give a flavour of two different widely popular processors we introduce the Motorola 68K series and the Sun SPARC processors. The Motorola 68K processors are very popular in high performance embedded applications while the Sun SPARC processors are popular in very high end enterprise servers. We will compare them with the Intel x86 series which is known for its success in the desktop market.

### 10.1. MOTOROLA 68K PROCESSORS

Motorola 68K processors are very similar to Intel x86 series in their architecture and instruction set. The both are of the same era and added various features at the same time. The instructions are very similar however the difference in architecture evident from a programmer's point of view must be understood.

68K processors have 16 23bit general purpose registers named from A0-A7 and D0-D7. A0-A7 can hold addresses in indirect memory accesses. These can also be used as software stack pointers. Stack in 68K is not as rigid a structure as it is in x86. There is a 32bit program counter (PC) that holds the address of currently executing instruction. The 8bit condition code register (CCR) holds the X (Extend) N (Negative) Z (Zero) V (Overflow) C (Carry) flags. X is set to C for extended operations (addition, subtraction, or shifting).

Motorola processors allow bit addressing, that is a specific bit in a byte or a bit field, i.e. a number of bits can be directly accessed. This is a very useful feature especially in control applications. Other data types include byte, word, long word, and quad word. A special MOVE16 instruction also accepts a 16byte block.

68K allows indirect memory access using any A register. A special memory access allows postincrement or predecrement as part of memory access. These forms are written as (An), (An)+, and -(An). Other forms allow addressing with another register as index and with constant displacement. Using one of the A registers as the stack pointer and using the post increment and pre decrement forms of addressing, stack is implemented. immediates can also be given as arguments and are preceded with a hash sign (#). Addressing is indicated with parenthesis instead of brackets.

68K has no segmentation; it however has a paged memory model. It used the big endian format in contrast to the little endian used by the Intel processors. It has varying instruction lengths from 1-11 words. It has a decrementing stack just like the Intel one. The format of instructions is "operation source, destination" which is different from the Intel order of operands. Some instructions from various instruction groups are given below.

#### Data Movement

```
EXG D0, D2
```

```
MOVE.B (A1), (A2)
```

```

MOVEA (2222).L, A4
MOVEQ #12, D7
Arithmetic
ADD D7, (A4)
CLR (A3) (set to zero)
CMP (A2), D1
ASL, ASR, LSL, LSR, ROR, ROL, ROXL, ROXR (shift operations)
Program Control
BRA label
JMP (A3)
BSR label (CALL)
JSR (A2) (indirect call)
RTD #4 (RET N)
Conditional Branch
BCC (branch if carry clear)
BLS (branch if Lower or Same)
BLT (branch if Less Than)
BEQ (branch if Equal)
BVC (branch if Overflow clear)

```

## 10.2. SUN SPARC PROCESSOR

The Sun SPARC is a very popular processing belonging to the RISC (reduced instruction set computer) family of processors. RISC processors originally named because of the very few rudimentary instructions they provided, are now providing almost as many instruction as CISC (complex instruction set computer). However some properties like a fixed instruction size and single clock execution for most instructions are there.

SPARC stands for Scalable Processor ARChitecture. SPARC is a 64bit processor. Its byte order is user settable and even on a per program basis. So one program may be using little endian byte order and another may be using big endian at the same time. Data types include byte, Halfword, Word (32bit), and Double Word (64bits) and Quadword. It has a fixed 32bit instruction size. It has a concept of ASI (Address Space Identifier); an 8bit number that works similar to a segment.

There are 8 global registers and 8 alternate global registers. One of them is active at a time and accessible as g0-g7. Apart from that it has 8 in registers (i0-i7), 8 local registers (l0-l7), and 8 out registers (o0-o7). All registers are 64bit in size. The global registers can also be called r0-r7, in registers as r8-r15, local registers as r16-r23, and out registers as r24-r31.

SPARC introduces a concept of register window. One window is 24 registers and the active window is pointed to by a special register called Current Window Pointer (CWP). The actual number of registers in the processor is in hundreds not restricted by the architecture definition. Two instruction SAVE and RESTORE move this register window forward and backward by 16 registers. Therefore one SAVE instruction makes the out register the in registers and brings in new local and out registers. A RESTORE instruction makes the in registers out registers and restores the old local and in registers. This way parameters passing and returning can be totally done in registers and there is no need to save and restore registers inside subroutines.

The register o6 is conventionally used as the stack pointer. Return address is stored in o7 by the CALL instruction. The register g0 (r0) is always 0 so loading 0 in a register is made easy. SPARC is a totally register based architecture, or it is called a load-store architecture where memory access is only allowed in data movement instruction. Rest of the operations must be done on registers.

SPARC instructions have two sources and a distinct destination. This allows more flexibility in writing programs. Some examples of instructions of this processor follow.

**Data Movement**

|               |                      |
|---------------|----------------------|
| LDSB [rn], rn | (load signed byte)   |
| LDUW [rn], rn | (load unsigned word) |
| STH [rn], rn  | (store half word)    |

**Arithmetic**

|                        |            |
|------------------------|------------|
| source1 = rn           |            |
| source2 = rn or simml3 |            |
| dest = rn              |            |
| ADD r2, r3, r4         |            |
| SUB r2, 4000, r5       |            |
| SLL, SRA, SRL          | (shifting) |
| AND, OR, XOR           | (logical)  |

**Program Control**

|           |                           |
|-----------|---------------------------|
| CALL      | (direct call)             |
| JMPL      | (register indirect)       |
| RET       |                           |
| SAVE      |                           |
| RESTORE   |                           |
| BA label  | (Branch Always)           |
| BE label  | (branch if equal)         |
| BCC label | (branch if carry clear)   |
| BLE label | (branch if less or equal) |
| BVS label | (branch if overflow set)  |